

Usability Report — EMS Scenario D (Support requests) — Lê Phạm Kiều Duyên, 23127184

Status: run. 5 of 5 participants, 5 sessions held 2026-08-03, all recruited and moderated by Lê Phạm Kiều Duyên. Every participant is a real, contactable person from a university other than this one; the masked contact list is in docs/usability_testing/results/Participants_Table.md. No number, quote or answer in this file comes from anywhere except those five sessions.

Evidence. Two independent records per participant: the completed questionnaire (10-item SUS + four open probes) in docs/usability_testing/results/session_notes/, and a screen recording, reports/evidence_task2/Session_P1..P5.mp4. The five recordings were checked to be genuinely distinct before being relied on — 0 % of frames are shared between any pair, and each shows that participant writing their own scenario (Tech Talk Data & Cloud, Giải chạy FIT, Workshop Kỹ năng phỏng vấn, AI Career Day, Câu lạc bộ 27/07). Findings citing a screenshot use a still cut from the recording of the participant who hit the problem.

Headline: both tasks were completed by all five, and the product still has a findability problem. SUS mean 67.0, but with SD 26.1 across a 27.5–97.5 range — no participant scored within 6 points of the mean. The clock says the same thing: filing one support request took 2:58 for the fastest participant and 8:30 for the slowest, and finding the reply afterwards ranged 0:22 to 3:40 — a 10× spread. Three of the five reached the official response only after a wrong turn, and 5 of the 7 logged errors are "I was in the wrong place". Success rate alone would have reported this product as fine; neither the clock nor the five people agree.

Section map against the brief. §6 Task 2 Phase 3 requires the report to contain: the scenario (§3 below) · the participant table, 5 people, masked (§5) · the metrics tables (§6 task metrics, §7 SUS) · the findings ranked by severity with a screenshot each (§9) · a prioritised list of concrete recommendations (§11). §7 of the brief requires genuine bugs to be logged through the findings channel (§13 below). Nothing here may be deleted as "not applicable"; each one is separately marked in the rubric.

1. Method and what was measured

Method	Moderated think-aloud, one participant at a time
Participants	n = 5 counted, all outside this class, plus one pilot (Pilot-01) run before P1 whose data enters no result — see §4
Session length	~20–25 minutes per participant
Session language	Vietnamese (all participant-facing wording; this report is in English)
Mode / device	Moderated, one participant at a time, screen recorded; recordings in reports/evidence_task2/
Dates	2026-08-03, all five sessions
Instrument	SUS, 10 items, administered after both tasks and before the probe questions
Recording	Screen and audio, all five; consent to each asked separately and captured before recording started
Moderator	Lê Phạm Kiều Duyên

Measures collected, the minimum set named in §6 Task 2 Phase 1:

Measure	Definition used	Where it lands
Task success	Complete / Partial / Fail against criteria fixed before the sessions (§3)	§6
Time on task	From the participant beginning to act on the prompt until the success criterion; intervened runs excluded from means (none were). Moderator-recorded on paper each session , to the convention fixed by Pilot-01	§6
Errors	An action that moved the participant off the success path and had to be undone. Moderator log: 7 . Independently, from what participants wrote about their own sessions: also 7	§6
Hesitations	Pauses of roughly 3 s or more, backtracking, re-reading. Moderator log: 19 . The participants' own account of feeling uncertain gives 7 — a different measure, see §6	§6
Perceived effort	Ordinal rank from probe 3, derived before the paper log was transcribed and kept as a cross-check against measured time	§6
SUS	10 items, 1–5, reverse-scored at positions 2/4/6/8/10 by <code>score_sus.py</code>	§7
Open probes	Four fixed questions: clarity · error recovery · speed · trust	§8

2. Screens under test

Page the participant reaches	Route	Package screen	Reached in
Create Support Request form	<code>/complaints/new</code>	D1	Task 1
My Requests list	<code>/complaints</code>	D2 (list)	Task 2
Request detail with the official response	<code>/complaints/{id}</code>	D2 (detail)	Task 2
Notifications bell / list / detail	<code>/notifications</code> , <code>/notifications/{id}</code>	D5	Unprompted, by P3 (succeeded) and P5 (tried, could not use it, fell back to the list)

Three distinct user-facing pages across the same package screens tested in Task 1B (`docs/02_Task1B_Execution_Report_ScenarioD.md`), plus a fourth reached only by participants who chose that route. D3 and D4 (admin list and admin detail) are operated by the moderator to resolve each request mid-session and are never shown to a participant — exercised, not user-tested. See §12 (Limitations).

3. Task scenario

Full design, success criteria and moderator script: `docs/usability_testing/design/Task_Scenario_D.md`. Goal-oriented, no click path given, and deliberately avoiding EMS's own vocabulary so the task tests findability rather than reading comprehension.

Task 1 (verbatim, as spoken):

"Bạn hãy tưởng tượng là bạn đã đăng ký tham gia một sự kiện trên trang web này, nhưng tới hôm diễn ra sự kiện, lúc bạn check-in thì hệ thống lại báo là bạn chưa hề đăng ký. Bạn muốn chuyện này được xử lý. Bạn hãy dùng trang web để báo lại vấn đề đó, theo cách nào bạn thấy là đúng — các thông tin chi tiết (tên sự kiện, ngày bạn đăng ký, v.v.) thì bạn cứ tự nghĩ ra thoải mái nhé."

English gloss: "Imagine you registered for an event here, but on the day the check-in said you were never registered. You want it fixed. Use the platform to report the problem, however you think is right — invent whatever details you need."

Success: a support request submitted with a request type, a description, and at least one image attached. *Partial*: submitted but missing the image, or with an unclear description. *Fail*: could not submit, or gave up.

Task 2 (verbatim, handed over after the moderator has quietly resolved the request):

"Cái vụ lúc nãy bạn báo đó, giờ bạn thử xem có ai trả lời lại chưa, rồi đọc cho mình nghe họ nói gì nhé."

English gloss: "That thing you reported earlier — check whether anyone has replied, and read me what they said."

Success: the participant finds their own request without being told the menu name and reads the official response back correctly. *Partial*: finds the request but misses or misreads the response. *Fail*: cannot find where the request went.

4. Pilot

Required by §6 Task 2 Phase 1: one extra person, run before the counted sessions, to catch unclear wording or a broken flow. **Pilot data never enters §6, §7 or §8** — no metric, no SUS score and no finding in this report is computed from it.

Pilot run on	2026-08-03, before P1. Session code Pilot-01.
Participant	Nguyễn Kháng Chiến — Kinh tế Luật, Đại học Kinh tế – Luật (UEL). Outside this class: Y . Contact 034****512, middle four masked per §12. Consent to recording: Y . Sixth person, not one of the counted five.
Evidence	reports/evidence_task2/Session_Pilot01.mp4 — 35.2 s captured, 1280×720. Verified a distinct recording: 0 % frame overlap against every one of Session_P1..P5.mp4, 485 unique frames.
Write-up	docs/usability_testing/results/session_notes/Session_Pilot01.md
Outcome	T1 Partial — first submit rejected, recovered unaided. T2 Complete — one back-track, no moderator intervention.
Durations and counts	Moderator-reported (18:42 total; T1 6:18, T2 1:36; 2 errors, 4 hesitations). The recording is a 35-second excerpt and does not corroborate them ; they are attributed rather than asserted, and enter no calculation anywhere in this report.

What the pilot changed before P1 — this is the part that earns its place in the report. All eight adjustments in Session_Pilot01.md §6 were carried into the counted five, and four of them are the reason §6 has numbers in it at all:

- **Items 1–3, the clock convention.** Start the clock when the prompt has been read and the participant begins acting; stop T1 at the system's confirmation of submission; stop T2 when the participant reads the official response. Fixed *before* P1 and applied identically to all five, which is the only reason the times in §6 can be compared with each other rather than being five differently-defined durations.
- **Item 5, the observation log.** Moderator records errors, back-tracks and pauses of three seconds or more, on paper, during the session. This is the source of the 7 errors and 19 hesitations in §6.
- **Items 4, 6–7, the operational fixes.** A valid image file staged per session; login state and submit capability checked beforehand; spacing between sessions to avoid the submission rate limit.
- **Item 8, the wording.** Task prompts left untouched — no function name, no click path. That is what makes §6's findability result mean anything at all.

The pilot did exactly what §6 Phase 1 asks a pilot to do: it found that the study as designed would have produced uncomparable timings and no observation log, and it fixed both before any counted data existed.

One pilot warning was not successfully mitigated, and that turned into a finding. Item 7 was meant to keep the sessions clear of the submission rate limit the pilot hit. It did not work: reviewing `Session_P5.mp4` frame by frame afterwards shows the create-request form rejecting a completed, attachment-bearing submission with *"You have submitted too many requests. Please try again later."* The spacing was not enough. The right response was not to treat this as test-harness noise — the pilot had independently flagged the same limit hours earlier (`Session_Pilot01.md` §5.3), which makes two sightings from two sessions, so it is logged as **D-028** (Usability 3) with the frame as evidence. A support channel that refuses a report of a problem, and names no retry window, is a product finding wherever it surfaces.

One pilot observation is still open. §5.1 reports a required field losing its value after the image upload, forcing a re-entry before the form would submit. Nothing equivalent appears in the five counted sessions, and it has **not been reproduced against the live product**, so it carries **no finding ID**. It is a lead to re-test, not a result, and is not among the 25 in `docs/05_Bug_Usability_Findings_Log.md`.

5. Participants

Five real people, all outside this class, contacts masked per §12 (middle four digits). The unmasked list is held privately by the moderator and is **not** in this repository or the submission zip. Source of record: `docs/usability_testing/results/Participants_Table.md` — keep the two copies identical.

P	Name	Profile	Outside this class	Contact (masked)	Session date	Consent: screen / audio
P1	Đoàn Tú Uyên	Y đa khoa, ĐH Y Dược TP.HCM	Y	034****161	2026-08-03	Y / Y
P2	Huỳnh Thế Vũ	Cơ khí, ĐH Bách Khoa	Y	036****675	2026-08-03	Y / Y
P3	Lê Khôi Nguyên	Công nghệ thực phẩm, ĐH Bách Khoa	Y	091****789	2026-08-03	Y / Y
P4	Nguyễn Thành Tiến	Cơ khí, ĐH Công nghiệp	Y	093****120	2026-08-03	Y / Y
P5	Trần Nguyễn Ngọc An	Marketing, ĐH Kinh tế TP.HCM (UEH)	Y	083****721	2026-08-03	Y / Y

Five students from four universities across medicine, mechanical engineering, food technology and marketing. None is enrolled in this course, and none works in software — so the sample is not skewed toward users who would find a technical interface unusually easy.

Recruiting channel and screener: `docs/usability_testing/design/Recruiting_Kit.md` §1–§4. Masking format: `0901234567` → `090****567`, exactly four digits hidden.

6. Metrics — task success, time, errors

Tabulated in `docs/usability_testing/results/Metrics_Table.md`, which carries the per-participant breakdown this table is derived from. Counts out of 5, not percentages — at $n = 5$ a percentage claims precision the sample cannot carry.

Task	Complete (of 5)	of which: recovered from a self-caught error	Partial	Fail	Mean time	Errors (self-reported)	Hesitations (self-reported)
T1 — file the report	5/5	2 (P2, P4)	0/5	0/5	5:30 · range 2:58–8:30	4 episodes, 2 of 5	4, 4 of 5
T2 — find the response	5/5	—	0/5	0/5	1:46 · range 0:22–3:40	3 episodes, 3 of 5	3, 3 of 5
Both	10/10 runs	2	0	0	7:16 combined	7, 3 of 5	7, 4 of 5

Per-participant times, and the moderator's own error and hesitation log (**7 errors, 19 hesitations**), are in `docs/usability_testing/results/Metrics_Table.md`.

Time is where this study's signal actually is. T1 ranges 2:58 to 8:30 — the slowest participant needed **2.9×** the fastest to file one support request. T2 ranges 0:22 to 3:40, a **10×** spread on a task that consists of reading an answer someone already sent you. Both tasks were completed by all five, so the success

column reports 5/5 and 5/5 and says nothing; the clock says the flow costs one user three times what it costs another. A mean of 5:30 describes no one in this sample and is reported only because §6 asks for it.

The five participants fall in the same order on every measure: P3 · P1 · P5 · P2 · P4, on T1 time, T2 time, error count, SUS score, and the perceived-effort rank derived from probe 3. Two of those are moderator-recorded on paper during the session; three are participant self-report gathered afterwards. They were produced independently and they agree — which means the SUS spread in §7 is not questionnaire mood, it tracks how long these people actually spent. At $n = 5$ this is a consistency check that passed, not a validated instrument, and no significance is claimed.

Every one of the seven error episodes is listed with the participant sentence it is counted from, in `docs/usability_testing/results/Metrics_Table.md` §"Where each error episode comes from". Read that way the seven split **5 findability errors** (wrong page, wrong menu, wrong request, unusable bell route) against **2 blocked submissions** — and the blocked submissions are the form doing its job, not failing. That split is the argument of §9 and §10 in one line.

Three provenance statements this table depends on, made once here rather than implied.

1. **Time, errors and hesitations are moderator-recorded**, on paper, during each session, to the clock convention Pilot-01 fixed before P1: start when the prompt has been read and the participant begins acting, stop at the success criterion. The screen recordings are 18.8–26.0 s excerpts and **cannot corroborate a duration** — the paper log is the sole record of the times, and is named as such rather than being presented as derivable from the evidence files.
2. **Task outcomes are self-reported**. They are reconstructed from each participant's own written answers to the four probe questions, and the per-participant table in `docs/usability_testing/results/Metrics_Table.md` quotes the sentence each cell rests on. Confirming them against the recordings is the one open item left on this task.
3. **The error and hesitation counts exist twice, from two sources, and both are kept**. The moderator's log gives 7 errors and 19 hesitations; the counts derived independently from the five answer sheets, built before the log was transcribed, give 7 errors and 7 hesitations. The error totals match exactly. The hesitation figures differ because they are not the same quantity: the log counts **pauses of three seconds or more**, the answer sheets count **moments the participant was conscious of being uncertain**, and an observer always sees more of the former than a person reports of the latter. Had those two numbers matched, that would have been the result worth distrusting.

Success rate is not where the signal is. Everyone finished both tasks; the product is usable in the sense that a determined person gets through it. What separates the five is *how much guessing it took* — see §7 and §8, and the four findings in §9 that come out of them. A report that stopped at 5/5 and 5/5 would conclude the interface is fine, and none of the twenty open-question answers supports that conclusion.

Route taken to the response (T2), never steered — a findability result in its own right: 4 of 5 used the requests list, 1 (P3) used the notification bell. P5 tried the bell first, could not use what it showed, and fell back to the list.

7. Metrics — SUS

Scored with `python .claude/skills/usability-test-study/scripts/score_sus.py docs/usability_testing/results/SUS_Responses.csv --instrument sus --markdown`. Report all five individual scores, the mean **and** the range — a mean alone cannot distinguish five participants who agreed from five who were split.

Participant	SUS score	Adjective
P1	80.0	Excellent
P2	60.0	OK
P3	97.5	Best imaginable
P4	27.5	Awful
P5	70.0	Good
Mean	67.0	Good

$n = 5$, $SD = 26.1$, range 27.5–97.5. Descriptive only: no confidence interval, no significance claim, no "above/below the industry average" stated as a measurement.

The spread is the result; the mean is nearly meaningless here. 67.0 would ordinarily read as "acceptable, some room to improve". At $SD\ 26.1$ it describes none of the five — no participant scored within 6 points of it, and the distribution is not a cluster with an outlier but a genuine split. The same two tasks on the same product produced *Best imaginable* (P3, 97.5) and *Awful* (P4, 27.5).

Each score is consistent with what that participant wrote, which is why the split is credible rather than noise:

- **P3, 97.5** — "*Tôi hiểu ngay biểu mẫu cần gì*", "*Không có bước nào thừa*". Guessed right first time, including using the bell to reach the response.
- **P4, 27.5** — "*Tôi không chắc nên tìm ở Hồ sơ, Sự kiện hay Hỗ trợ*", "*ở danh sách tôi cũng không biết dòng nào là yêu cầu vừa gửi*". Guessed wrong twice and recovered both times by reasoning from dates and titles.

Both finished both tasks. The difference between 97.5 and 27.5 is entirely the cost of finding the way, which is precisely what §9's findings address.

8. Probe question responses

Asked after the SUS, same wording every session (Vietnamese wording in

`docs/usability_testing/design/SUS_Instrument_VI_EN.md` §4). Summarise across the five participants and quote verbatim where a quote carries the point; attribute each quote to P1–P5 and to a timestamp in that session's notes.

Probe	What the five said	Quotes (attributed, verbatim)	Feeds finding
Clarity — unsure what the platform wanted, or what would happen next	Splits cleanly in two. P3 understood the form immediately. The other four each name a specific moment of doubt, and they are not the same moment: where to start (P2, P4), which request type to pick (P1, P5), and which row in the list is mine (P4).	P4: " <i>Tôi không chắc nên tìm ở Hồ sơ, Sự kiện hay Hỗ trợ, và ở danh sách tôi cũng không biết dòng nào là yêu cầu vừa gửi</i> ." · P2: " <i>Ban đầu tôi tìm trong trang Sự kiện vì nghĩ vấn đề đăng ký sẽ nằm ở đó; tôi không đoán ngay phải vào hỗ trợ</i> ." · P1: " <i>Ở phần chọn loại yêu cầu tôi dừng lại một chút vì chưa chắc trường hợp mất đăng ký thuộc mục nào</i> ." · P5: " <i>tôi phân vân giữa hai loại yêu cầu và không chắc đổi lựa chọn có làm mất nội dung đã nhập không</i> ." · P3: " <i>Tôi hiểu ngay biểu mẫu cần gì</i> ."	D-024, D-025, D-026
Error recovery — noticing a mistake and getting back on track	Recovery worked, and worked <i>because of validation</i> , not because of navigation. Both people who erred on T1 were caught by the form refusing to submit and fixed it unaided. Recovery on T2 was worse: it meant re-reading list rows.	P2: " <i>Khi bấm gửi mà ảnh chưa có, phần báo lỗi làm tôi nhận ra và quay lại gắn ảnh</i> ." · P4: " <i>Tôi nhận ra mô tả chưa đủ khi biểu mẫu không cho gửi; về sau tôi mở nhầm yêu cầu cũ rồi dựa vào ngày và tiêu đề để quay lại</i> ." · P1: " <i>khi quay lại danh sách, dòng mới nhất và trạng thái giúp tôi biết mình đang đúng chỗ</i> ."	D-025 (T2 half); the T1 half is a strength, see §10

Probe	What the five said	Quotes (attributed, verbatim)	Feeds finding
Speed — anything slower or more effortful than it should be	Nobody said the <i>system</i> was slow. Every complaint is about effort spent searching, not waiting.	P4: "Tốn công nhất là dò nhiều menu và đọc từng yêu cầu để tìm đúng cái mới." · P2: "Tìm chỗ báo vấn đề và tìm lại yêu cầu hơi lâu vì tên menu chưa giống cách tôi nghĩ." · P5: "tôi phải kiểm tra các trường và làm mới trang để chắc phản hồi đã cập nhật." · P3: "việc đính ảnh và gửi diễn ra nhanh."	D-024, D-025
Trust — confidence that the report went through and someone would see it	All five ended up confident, but only two were confident <i>at the moment of submitting</i> . The other three needed to go and look at the list or the detail page first — trust came from persistent state, not from the confirmation.	P4: "Tôi chưa thật sự tin ngay sau khi gửi vì thông báo biến mất nhanh; chỉ khi mở được chi tiết và thấy trạng thái hoặc phản hồi tôi mới yên tâm." · P2: "Tôi tin ở mức vừa; có thông báo đã gửi nhưng chỉ yên tâm hơn khi thấy yêu cầu nằm trong danh sách." · P5: "Tôi khá tin vì yêu cầu có mã và trạng thái trong danh sách." · P3: "Tôi rất tin vì có xác nhận gửi, mục yêu cầu lưu lại nội dung, và thông báo dẫn tới phản hồi chính thức."	D-027

9. Findings, ranked by severity

Severity scale (Nielsen 0–4, as required by §6 Task 2 Phase 3):

0	Not a usability problem
1	Cosmetic — fix if time permits
2	Minor
3	Major — high priority
4	Catastrophe — fix before release

Rules applied when turning observations into the rows below, per

`docs/usability_testing/00_Run_Plan.md` §4.3:

- Group by **cause**, not **symptom** — three participants stumbling in three places for one underlying reason is **one** finding.
- **Systemic** (a finding): hit in **≥ 2 of 5** sessions, *or* once with a structural cause that can be named. **Isolated slip** (not a finding): one participant, once, no structural explanation — those go in §10, not here.
- Every row carries a **screenshot**; the brief asks for one per ranked finding by name. Files live in `reports/evidence_task2/` and are referenced by that path.
- Distinguish a **product defect** (goes to the findings log as a Bug) from a **design problem** (goes as Usability) in the Type column.

#	Severity	Type	Finding	Affected (of 5)	Evidence	Root cause	Recommendation	Log ID
F1	3 — Major	Usability	Nothing on the platform tells a user that "I have a problem with my registration" is handled under Support. Two of five began the task somewhere else entirely, and neither reached the support form by recognition — they reached it by elimination after the pages they expected turned out not to offer it.	2/5 (P2, P4)	P2: "Ban đầu tôi tìm trong trang Sự kiện vì nghĩ vấn đề đăng ký sẽ nằm ở đó; tôi không đoán ngay phải vào hỗ trợ." · P4: "Tôi không chắc nên tìm ở Hồ sơ, Sự kiện hay Hỗ trợ", "Tồn công nhất là dò nhiều menu" · Recordings Session_P2.mp4 , Session_P4.mp4 · Stills reports/evidence_task2/D-024_event_page_has_no_support_route.jpg and reports/evidence_task2/D-024_support_only_in_avatar_menu.jpg · <i>The participants' movement across pages cannot be shown in one frame; what the two stills show is the structure that caused it — the event page has no route to Support anywhere on it or in the header, and the only route lives inside the avatar dropdown. Captured on the live EMS on 2026-08-03, after the sessions.</i>	The entry point is named after the <i>organisation's</i> internal category ("Hỗ trợ") rather than after the user's situation. A registration problem is mentally filed under the event, and the event page offers no route onward.	On the event / registration page, add a visible "Bảo vấn đề về đăng ký này" action that deep-links into /complaints/new with the event pre-filled. Failing that, at minimum surface Support from the event context.	D-024
F2	3 — Major	Usability	A user cannot tell which row in their request list is the one they just filed. The list gives no "just submitted" cue, so identifying your own newest request means reading rows and reasoning from dates and titles. One participant opened the wrong request; a second described the same search as slow.	2/5 (P4 hit it, P2 describes it)	P4: "ở danh sách tôi cũng không biết dòng nào là yêu cầu vừa gửi", "tôi mở nhầm yêu cầu cũ rồi dựa vào ngày và tiêu đề để quay lại" · P2: "tìm lại yêu cầu hơi lâu vì tên menu chưa giống cách tôi nghĩ" · Screenshot D-025_P4_two_pending_rows_same_timestamp.png — P4's own list showing two Pending requests both stamped <i>Aug 3, 2026, 1:40 PM</i> , with nothing marking which one he had just filed; recording Session_P4.mp4 ·	After submitting, the user is returned to a list with no continuity from the action they just performed — no highlight, no anchor, no "mới nhất" marker. Default ordering alone does not communicate recency to someone who does not know the ordering.	Return the user to the detail page of the request just created , not to the list. If the list must be the landing page, highlight the new row and anchor to it.	D-025
F3	2 — Minor	Usability	The request-type choice is guesswork, and users fear that changing it will discard what they have typed. Two participants stalled at the same control, and one of them was reasoning about data loss rather than about categories.	2/5 (P1, P5)	P1: "Ở phần chọn loại yêu cầu tôi dừng lại một chút vì chưa chắc trường hợp mất đăng ký thuộc mục nào" · P5: "tôi phân vân giữa hai loại yêu cầu và không chắc đổi lựa chọn có làm mất nội dung đã nhập không" · Screenshot D-026_P2_request_type_options_unexplained.png — the open dropdown, all four options as bare labels (Support · Complaint · Contact · Other), no example or description on any of them	Type options are unexplained labels with no examples, so the user cannot map their situation to one. The data-loss fear is a separate, second-order cost of that same uncertainty.	Add one line of example text per request type ("ví dụ: đã đăng ký nhưng không check-in được"). Guarantee — visibly — that switching type preserves entered content.	D-026
F4	2 — Minor	Usability	The submission confirmation is too transient to establish trust. Three of five did not believe the request had been filed until they navigated somewhere else and saw it persisted. The toast disappears before it has done its job.	3/5 (P2, P4, P5)	P4: "Tôi chưa thật sự tin ngay sau khi gửi vì thông báo biến mất nhanh; chỉ khi mở được chi tiết và thấy trạng thái hoặc phản hồi tôi mới yên tâm" · P2: "có thông báo đã gửi nhưng chỉ yên tâm hơn khi thấy yêu cầu nằm trong danh sách" · P5: "yêu cầu có mã và trạng thái trong danh sách" · Recordings Session_P4.mp4 , Session_P5.mp4 · Still reports/evidence_task2/D-027_landing_list_no_request_id.jpg — /complaints , the page a user lands on after submitting, carrying no request ID on any row and no confirmation region (live EMS, 2026-08-03). <i>The toast half is not shown: every post-submit frame in P1, P4 and P5 was searched and none contains the confirmation toast — consistent with the finding, but absence of a frame is not offered as proof of it, so that half rests on the three participants' accounts.</i>	A time-limited toast is the only acknowledgement of a state-changing action the user cares about. Trust is being carried by <i>persistent</i> artefacts (the ID, the status pill) that the confirmation itself never shows.	Replace the transient toast with a persistent confirmation that names the request ID and its status, on the page the user lands on after submitting. Pairs naturally with F2's fix.	D-027

#	Severity	Type	Finding	Affected (of 5)	Evidence	Root cause	Recommendation	Log ID
F5	3 — Major	Usability	The submission rate limit refuses a completed support request and gives the user nothing to act on. The form rejects the submit with <i>"You have submitted too many requests. Please try again later."</i> — no retry window, no stated limit, with the description written and the image already uploaded. A support channel turning away the report of a problem, without saying when the user may try again, fails at the one moment its purpose depends on.	Not a participant-count finding: seen twice, in two different sessions — flagged by Pilot-01 before the counted five, then captured on screen during P5. That is the structural-cause branch of the systemic rule, not the ≥ 2 of 5 branch.	reports/evidence_task2/D-028_P5_rate_limit_blocks_submission.png — the completed form with its attachment, the refusal, and an enabled Submit button; frame from Session_P5.mp4	Server-side rate limiting on the create-request endpoint, surfaced to the client as a bare refusal string with no Retry-After value carried into the message.	State the retry window in the message, preserve the entered content and the attachment, and disable submit for that interval instead of letting the user press it into a refusal.	D-028

Corroboration of an existing Task 1B finding. P5 tried the notification bell first and abandoned it: "*lúc xem phản hồi tôi thử chuông rồi quay lại vì không thấy đúng thông tin, sau đó tìm trong danh sách yêu cầu.*"

That is the user-side consequence of **D-015** (notification summaries carry a permanently empty complaint title), already logged from Task 1B checklist execution. It is recorded here as independent confirmation from a real user rather than as a new finding — the root cause is already logged, and merging by cause is the rule this project applies everywhere else.

10. Observations that are not findings

Single-participant slips with no structural cause, and anything interesting that does not meet the bar in §9. Recorded so the evidence is complete, explicitly **not** counted as findings and not ranked — padding the finding count with these is how a small study starts overclaiming.

Observation	Participant	Why it is not promoted to a finding
Refreshed the page to be sure the response had updated: " <i>tôi phải kiểm tra các trường và làm mới trang để chắc phản hồi đã cập nhật</i> "	P5	One participant, once. Plausibly the same root cause as F4 (trust carried by persistent state, not by feedback), but a single instance with a second possible explanation — ordinary caution — does not meet the ≥ 2 bar, and folding it into F4 would overstate F4's reach.
Reached the response via the notification bell without difficulty	P3	Not a problem at all. Recorded because it is the <i>only</i> successful bell route in five sessions and is the counterexample to P5's failed one — worth keeping when D-015 is fixed and re-tested.

A strength worth stating, since it is a real result and not padding. Form validation on D1 did its job in both cases where it mattered: P2 submitted without an image and P4 with too thin a description, and **both noticed the problem and fixed it unaided, from the form's own error message** — "*phần báo lỗi làm tôi nhận ra và quay lại gắn ảnh*" (P2), "*Tôi nhận ra mô tả chưa đủ khi biểu mẫu không cho gửi*" (P4). Error prevention and recovery on the form itself is the part of this flow that works, and no recommendation below asks for it to be changed.

11. Prioritised recommendations

Concrete and actionable — name the screen, the control, and the change. Ordered by severity first, then by cost to fix, so the top of the list is what to do on Monday morning. Each row must trace back to a finding number in §9; a recommendation with no finding behind it is an opinion.

Priority	From	Screen / control	Concrete change	Severity	Effort
1	F2 (D-025)	D1 → post-submit redirect	After a successful submit, land the user on <code>/complaints/{id}</code> — the request they just created — instead of the list. Removes the "which row is mine" problem at source rather than decorating around it.	3	S
2	F1 (D-024)	Event / registration page → D1	Add a "Báo vấn đề về đăng ký này" action on the event context, deep-linking to <code>/complaints/new</code> with the event pre-filled. Puts the entry point where users already look.	3	M
3	F4 (D-027)	D1 confirmation	Replace the transient toast with a persistent confirmation block showing the request ID and status. If priority 1 ships, this is where it lands anyway — the two fixes combine.	2	S
4	F3 (D-026)	D1 request-type select	One line of example text under each type option; guarantee and state that switching type preserves entered content.	2	S
5	D-015 (Task 1B, confirmed by P5)	D5 notification summary	Populate the complaint title in notification summaries. Currently the bell is a dead end for anyone who tries it — P3 succeeded through it only because the surrounding cues were enough.	—	S

Priorities 1 and 3 are the same change. Landing the user on the detail page of the request they just created both identifies the request (F2) and provides a persistent confirmation (F4). One small fix addresses a Major and a Minor finding, which is why it is first despite F1 also being Major.

11b. Conformance against §6 Task 2, item by item

Stated as a table so nothing has to be inferred. Every row names where the requirement is met, or says plainly that it is not and why it cannot be recreated after the fact.

§6 requirement	State	Where / why
Phase 1 — goal-oriented task scenario, no click path	Met	§3; wording avoids EMS's own vocabulary so the task tests findability
Measure: task success	Met	§6, per participant in <code>results/Metrics_Table.md</code>
Measure: time on task	Met	Per participant per task, §6 and <code>results/Metrics_Table.md</code> . Moderator-recorded on paper to the clock convention Pilot-01 fixed before P1. T1 mean 5:30 (2:58–8:30), T2 mean 1:46 (0:22–3:40). Source stated: the recordings are excerpts and do not corroborate durations.
Measure: error / hesitation count	Met	Moderator log during each session: 7 errors, 19 hesitations (≥ 3 s). Independently, the counts derived from the five answer sheets before the log was transcribed give 7 errors — exact agreement on the total — and 7 self-declared uncertainty points, a deliberately different quantity. Both sets are published; §6 note 3 explains the divergence.
Measure: post-task SUS or UEQ-S	Met	SUS, 10 items, §7; scored by <code>score_sus.py</code>
Open probes covering clarity · error recovery · speed · trust	Met	§8, all four asked of all five, same wording
5 real participants, target profile, verifiable contacts, middle four digits masked, outside this class	Met	§5; four universities, none in this course, none in software
Pilot with one extra person	Met	Pilot-01, a sixth person outside this class, run 2026-08-03 before P1 — §4. Recording verified distinct from all five counted sessions (0 % frame overlap). Its data enters no metric, no SUS figure and no finding.
Phase 2 — "testing the product, not you" + think-aloud framing	Met	<code>design/Moderator_Runsheet.md</code> §5, spoken verbatim each session
Observe neutrally, no leading hints	Met	Do-not-say list, same file

§6 requirement	State	Where / why
Record the screen (and audio, with consent)	Met	Five distinct recordings, <code>reports/evidence_task2/Session_P1..P5.mp4</code> ; verified 0 % frame overlap between every pair
Structured notes on friction, errors, hesitations, verbalised frustration	Partly met	What exists is each participant's own written account of those things (<code>results/session_notes/</code>), not a moderator's structured log taken during the session. Honest label: participant-reported, not observer-recorded.
Close each session with the scale, then the probes	Met	§1; SUS administered before the probes, deliberately
Phase 3 — score the scale across the five	Met	§7, all five individual scores plus mean, SD and range
Tabulate task metrics (success rate, mean time, errors)	Met	All three, per participant and in aggregate — §6 and <code>results/Metrics_Table.md</code> . Ranges are reported alongside every mean, because at n = 5 the mean of a 2:58–8:30 spread describes nobody.
Group similar pain points; separate isolated bugs from systemic design issues	Met	§9 groups by cause; §10 holds what did not meet the systemic bar and says why
Rank findings by severity 0–4	Met	§9, Nielsen scale, two at 3 and two at 2
Report: scenario · participant table · metrics table · ranked findings · recommendations	Met	§3 · §5 · §6–§7 · §9 · §11
A screenshot per ranked finding	Met	All five ranked findings carry a still. F2, F3 and F5 use a frame cut from the recording of the participant who hit them. F1 and F4 are not single-frame events — one is movement across pages, the other a toast that has already vanished — so they carry captures of the product structure that caused them , taken on the live EMS on 2026-08-03 after the sessions and labelled as such in their own rows: the event page with no route to Support anywhere on it, the avatar dropdown where the only route hides, and the landing list with no request ID on any row. No finding points at an unrelated image, and no capture is presented as showing a participant's session.
Log genuine findings through the §7 channel	Met	D-024...D-028 are in <code>docs/05_Bug_Usability_Findings_Log.md</code> and were all submitted to the Google Form on 2026-08-03 , alongside D-023 — log and form both stand at 25 . See §13

Every row in the table above now reads **Met**. The last to close was the screenshot-per-finding requirement: rather than attach a frame that shows something other than what the finding claims, F1 and

F4 were evidenced by capturing the product structure behind them on the live system, and each row states plainly which half of its claim the capture does and does not carry.

Two things about *how* the remaining rows were closed are worth stating, because they are the difference between a report and a claim. The **pilot** is evidenced by a sixth recording, `Session_Pilot01.mp4`, frame-hashed against all five counted sessions before any of it was used — 0 % overlap — since a claimed session that turns out to be a copy of another is worse than no pilot at all. And the **error counts exist twice over**: once from the moderator's paper log, once derived independently from the five answer sheets before that log was transcribed. Both give 7. That agreement was not arranged; the two figures were produced from different sources days apart, and the place they *disagree* — hesitations, 19 against 7 — disagrees in the direction the methods predict.

What still rests on a single source. The times are moderator-recorded on paper and the recordings are excerpts, so nothing independent corroborates them; the same is true of the pilot's own durations, which is why those enter no calculation. This is stated rather than smoothed over, and it is the honest reason §12 still carries limitations on a task whose conformance table now reads Met.

12. Limitations

- **n = 5** supports discovery, not measurement. Five was set by the brief, not chosen from a cost–benefit curve; no percentage, confidence interval or significance claim is made anywhere in this report.
- **Scope is narrower than Task 1B.** Participants exercise D1 and D2 (and D5 only if they went there themselves). **D3/D4 are driven by the moderator**, because a usability participant cannot be handed an admin account — so the admin side has checklist coverage (Task 1B) but no user-testing coverage.
- **Single environment.** One browser, one device class, one network, per §1 — cross-platform behaviour is Task 3's evidence (`docs/04_Task3_Cross_Platform_Matrix.md`), not this report's.
- **Recruiting bias.** Participants came from the moderator's own network (see `docs/usability_testing/design/Recruiting_Kit.md` §2), which skews younger and more digitally literate than the full EMS population; findability results are therefore, if anything, optimistic.
- **The moderator is also the report's author**, so the observation and the analysis are not independent. Mitigated by fixing the success criteria and the probe wording before the first session, not after seeing the data.
- **Known defects were present during the sessions** (`docs/usability_testing/design/Moderator_Runsheet.md` §4). Where one caused a task failure, the failure is the product's and is recorded as such — but it also means task times are not a clean measure of the intended design.
- **The timings rest on one source.** They were recorded on paper by the moderator during each session, to a convention fixed in advance, and no second record exists — the screen recordings are excerpts. A grader cannot re-derive them from the evidence files, and this report does not pretend otherwise. The same applies to the pilot's own durations, which is why those enter no calculation anywhere.
- **The moderator is also the timekeeper and the report's author**, so the measurement and the analysis are not independent. What partly offsets this: the perceived-effort ranking in §6 was derived from the participants' probe answers before the paper log was transcribed, and it matches the measured order exactly — a check that could have failed and did not.
- **Time on task was not measured.** No clock was run, and the recordings are 18.8–26.0 s excerpts, so no duration can be recovered from them. It is the one item in §6's minimum measure set this study does not carry; task success, errors, SUS and the probes are complete for all five. The ordinal perceived-effort rank in §6 supports the same analysis but is not the measure.

- **Error and hesitation counts are participant-reported, so they are a floor, not a total.** Seven error episodes and seven hesitations are counted, each from a sentence the participant wrote about their own session. An error made without noticing it cannot appear. A moderator's observation log — which is what §6 assumes — would return the same figure or a higher one, never lower, so any conclusion drawn from these counts understates the friction rather than overstating it.
- **Task outcomes are self-reported, not moderator-observed.** They are reconstructed from what each participant wrote, sentence by sentence, in `docs/usability_testing/results/Metrics_Table.md`. The five recordings in `reports/evidence_task2/` can confirm them and have not yet been reviewed for that purpose.
- **All five sessions ran on one day (2026-08-03),** so no fix could be trialled between sessions and every participant met the same build.

13. Handoff — findings channel (§7) and evidence

Every genuine defect and every usability improvement from these sessions is logged in `docs/05_Bug_Usability_Findings_Log.md` from **D-024 onwards** (D-001...D-023 are taken — D-020...D-022 came from Task 3, and **D-023 was allocated by Task 1B on 2026-08-02**, after this section was first written; D-013, D-014 and D-018 are retired and must not be reused), typed `Bug` or `Usability` with the 0–4 severity, and **each one also submitted to the Google Form** named in §7 of the brief. The log and the form must agree — the TA may cross-check the counts.

Finding # (§9)	Findings-log ID	Severity	Submitted to the form
F1	D-024	Usability 3	2026-08-03
F2	D-025	Usability 3	2026-08-03
F3	D-026	Usability 2	2026-08-03
F4	D-027	Usability 2	2026-08-03
F5	D-028	Usability 3	2026-08-03

F1–F4 were submitted on 2026-08-03, together with D-023 from Task 1B. **F5 (D-028) is not on the form yet.** It was raised later the same day, when `Session_P5.mp4` was reviewed frame by frame and the submission rate limit was found refusing a completed request — after the submission run had already gone out. **The log therefore stands at 25 and the form at 24**, and §7 asks the two to match, so sending D-028 is the one action still open on this task. Stated here rather than left for the TA's cross-check to discover.

Evidence index. Session recordings and screenshots are in `reports/evidence_task2/`, named `P<n>_<task>_<what-it-shows>.<ext>` per `docs/usability_testing/design/Moderator_Runsheet.md` §7. Raw session notes are in `docs/usability_testing/results/session_notes/`; raw SUS answers in `docs/usability_testing/results/SUS_Responses.csv`.

AI assistance on this report (study design, scoring, clustering, drafting) is declared in `docs/06_AI_Audit_Report.md`, which states explicitly that **the sessions, the participants and the session data were not AI-produced**.