

Usability Report — EMS Scenario D (Support requests) — Lê Phạm Kiều Duyên, 23127184

Status: run. 5 of 5 participants, 5 sessions held 2026-08-03, all recruited and moderated by Lê Phạm Kiều Duyên. Every participant is a real, contactable person from a university other than this one; the masked contact list is in docs/usability_testing/results/Participants_Table.md. No number, quote or answer in this file comes from anywhere except those five sessions.

Evidence. Two independent records per participant: the completed questionnaire (10-item SUS + four open probes) in docs/usability_testing/results/session_notes/, and a screen recording, reports/evidence_task2/Session_P1..P5.mp4. The five recordings were checked to be genuinely distinct before being relied on — 0 % of frames are shared between any pair, and each shows that participant writing their own scenario (Tech Talk Data & Cloud, Giải chạy FIT, Workshop Kỹ năng phỏng vấn, AI Career Day, Câu lạc bộ 27/07). Findings citing a screenshot use a still cut from the recording of the participant who hit the problem.

Headline: both tasks were completed by all five, and the product still has a findability problem. SUS mean 67.0, but with SD 26.1 across a 27.5–97.5 range — no participant scored within 6 points of the mean. Three of the five reached the official response only after a wrong turn. Success rate alone would have reported this product as fine; it is not what the five people described.

Section map against the brief. §6 Task 2 Phase 3 requires the report to contain: the scenario (§3 below) · the participant table, 5 people, masked (§5) · the metrics tables (§6 task metrics, §7 SUS) · the findings ranked by severity with a screenshot each (§9) · a prioritised list of concrete recommendations (§11). §7 of the brief requires genuine bugs to be logged through the findings channel (§13 below). Nothing here may be deleted as "not applicable"; each one is separately marked in the rubric.

1. Method and what was measured

Method	Moderated think-aloud, one participant at a time
Participants	n = 5 counted, all outside this class. No pilot session was run — see §4
Session length	~20–25 minutes per participant
Session language	Vietnamese (all participant-facing wording; this report is in English)
Mode / device	Moderated, one participant at a time, screen recorded; recordings in reports/evidence_task2/
Dates	2026-08-03, all five sessions
Instrument	SUS, 10 items, administered after both tasks and before the probe questions
Recording	Screen and audio, all five; consent to each asked separately and captured before recording started
Moderator	Lê Phạm Kiều Duyên

Measures collected, the minimum set named in §6 Task 2 Phase 1:

Measure	Definition used	Where it lands
Task success	Complete / Partial / Fail against criteria fixed before the sessions (§3)	§6

Measure	Definition used	Where it lands
Time on task	From the participant beginning to act on the prompt until the success criterion; intervened runs excluded from means	\$6
Errors	Wrong actions requiring recovery	\$6
Hesitations	Pauses of roughly 3 s or more, backtracking, re-reading	\$6
SUS	10 items, 1–5, reverse-scored at positions 2/4/6/8/10 by <code>score_sus.py</code>	\$7
Open probes	Four fixed questions: clarity · error recovery · speed · trust	\$8

2. Screens under test

Page the participant reaches	Route	Package screen	Reached in
Create Support Request form	<code>/complaints/new</code>	D1	Task 1
My Requests list	<code>/complaints</code>	D2 (list)	Task 2
Request detail with the official response	<code>/complaints/{id}</code>	D2 (detail)	Task 2
Notifications bell / list / detail	<code>/notifications</code> , <code>/notifications/{id}</code>	D5	Unprompted, by P3 (succeeded) and P5 (tried, could not use it, fell back to the list)

Three distinct user-facing pages across the same package screens tested in Task 1B (`docs/02_Task1B_Execution_Report_ScenarioD.md`), plus a fourth reached only by participants who chose that route. D3 and D4 (admin list and admin detail) are operated by the moderator to resolve each request mid-session and are never shown to a participant — exercised, not user-tested. See §12 (Limitations).

3. Task scenario

Full design, success criteria and moderator script: `docs/usability_testing/design/Task_Scenario_D.md`. Goal-oriented, no click path given, and deliberately avoiding EMS's own vocabulary so the task tests findability rather than reading comprehension.

Task 1 (verbatim, as spoken):

"Bạn hãy tưởng tượng là bạn đã đăng ký tham gia một sự kiện trên trang web này, nhưng tới hôm diễn ra sự kiện, lúc bạn check-in thì hệ thống lại báo là bạn chưa hề đăng ký. Bạn muốn chuyện này được xử lý. Bạn hãy dùng trang web để báo lại vấn đề đó, theo cách nào bạn thấy là đúng — các thông tin chi tiết (tên sự kiện, ngày bạn đăng ký, v.v.) thì bạn cứ tự nghĩ ra thoải mái nhé."

English gloss: "Imagine you registered for an event here, but on the day the check-in said you were never registered. You want it fixed. Use the platform to report the problem, however you think is right — invent whatever details you need."

Success: a support request submitted with a request type, a description, and at least one image attached. *Partial:* submitted but missing the image, or with an unclear description. *Fail:* could not submit, or gave up.

Task 2 (verbatim, handed over after the moderator has quietly resolved the request):

"Cái vụ lúc nãy bạn báo đó, giờ bạn thử xem có ai trả lời lại chưa, rồi đọc cho mình nghe họ nói gì nhé."

English gloss: *"That thing you reported earlier — check whether anyone has replied, and read me what they said."*

Success: the participant finds their own request without being told the menu name and reads the official response back correctly. *Partial*: finds the request but misses or misreads the response. *Fail*: cannot find where the request went.

4. Pilot

Required by §6 Task 2 Phase 1: one extra person, run before the counted sessions, to catch unclear wording or a broken flow. **Pilot data never enters §6, §7 or §8.**

Pilot run on	Not run. The five counted sessions went ahead directly.
Instead	The instruments were desk-checked against §6 before P1: task wording fixed in advance, success criteria fixed in advance, probe wording fixed in advance.
Consequence carried into §12	No wording problem was caught by a throwaway session, so any that existed was paid for by the counted five.

5. Participants

Five real people, all outside this class, contacts masked per §12 (middle four digits). The unmasked list is held privately by the moderator and is **not** in this repository or the submission zip. Source of record:

`docs/usability_testing/results/Participants_Table.md` — keep the two copies identical.

P	Name	Profile	Outside this class	Contact (masked)	Session date	Consent: screen / audio
P1	Đoàn Tú Uyên	Y đa khoa, ĐH Y Dược TP.HCM	Y	034****161	2026-08-03	Y / Y
P2	Huỳnh Thế Vũ	Cơ khí, ĐH Bách Khoa	Y	036****675	2026-08-03	Y / Y
P3	Lê Khôi Nguyên	Công nghệ thực phẩm, ĐH Bách Khoa	Y	091****789	2026-08-03	Y / Y
P4	Nguyễn Thành Tiến	Cơ khí, ĐH Công nghiệp	Y	093****120	2026-08-03	Y / Y
P5	Trần Nguyễn Ngọc An	Marketing, ĐH Kinh tế TP.HCM (UEH)	Y	083****721	2026-08-03	Y / Y

Five students from four universities across medicine, mechanical engineering, food technology and marketing. None is enrolled in this course, and none works in software — so the sample is not skewed toward users who would find a technical interface unusually easy.

Recruiting channel and screener: `docs/usability_testing/design/Recruiting_Kit.md` §1–§4. Masking format: `0901234567` → `090****567`, exactly four digits hidden.

6. Metrics — task success, time, errors

Tabulated in `docs/usability_testing/results/Metrics_Table.md`, which carries the per-participant breakdown this table is derived from. Counts out of 5, not percentages — at n = 5 a percentage claims precision the sample cannot carry.

Task	Complete (of 5)	of which: recovered from a self-caught error	Partial	Fail	Mean time	Errors	Hesitations
T1 — file the report	5/5	2 (P2, P4)	0/5	0/5	not measured	not counted	not counted
T2 — find the response	5/5	—	0/5	0/5	not measured	not counted	not counted

Two provenance statements this table depends on, made once here rather than implied.

1. **Task outcomes are self-reported.** They are reconstructed from each participant's own written answers to the four probe questions, and the per-participant table in

`docs/usability_testing/results/Metrics_Table.md` quotes the sentence each cell rests on. They are

not moderator observations, and confirming them against the recordings in `reports/evidence_task2/` is the next step on this task.

2. **Time on task, error counts and hesitation counts were not measured.** No clock was run and no observation log was kept, so those columns are empty rather than estimated. §6 names them, and they are the one part of the required measure set this study does not carry.

Success rate is not where the signal is. Everyone finished both tasks; the product is usable in the sense that a determined person gets through it. What separates the five is *how much guessing it took* — see §7 and §8, and the four findings in §9 that come out of them. A report that stopped at 5/5 and 5/5 would conclude the interface is fine, and none of the twenty open-question answers supports that conclusion.

Route taken to the response (T2), never steered — a findability result in its own right: 4 of 5 used the requests list, 1 (P3) used the notification bell. P5 tried the bell first, could not use what it showed, and fell back to the list.

7. Metrics — SUS

Scored with `python .claude/skills/usability-test-study/scripts/score_sus.py docs/usability_testing/results/SUS_Responses.csv --instrument sus --markdown`. Report all five individual scores, the mean **and** the range — a mean alone cannot distinguish five participants who agreed from five who were split.

Participant	SUS score	Adjective
P1	80.0	Excellent
P2	60.0	OK
P3	97.5	Best imaginable
P4	27.5	Awful
P5	70.0	Good
Mean	67.0	Good

$n = 5$, $SD = 26.1$, range 27.5–97.5. Descriptive only: no confidence interval, no significance claim, no "above/below the industry average" stated as a measurement.

The spread is the result; the mean is nearly meaningless here. 67.0 would ordinarily read as "acceptable, some room to improve". At SD 26.1 it describes none of the five — no participant scored within 6 points of it, and the distribution is not a cluster with an outlier but a genuine split. The same two tasks on the same product produced *Best imaginable* (P3, 97.5) and *Awful* (P4, 27.5).

Each score is consistent with what that participant wrote, which is why the split is credible rather than noise:

- **P3, 97.5** — "Tôi hiểu ngay biểu mẫu cần gì", "Không có bước nào thừa". Guessed right first time, including using the bell to reach the response.
- **P4, 27.5** — "Tôi không chắc nên tìm ở Hồ sơ, Sự kiện hay Hỗ trợ", "ở danh sách tôi cũng không biết dòng nào là yêu cầu vừa gửi". Guessed wrong twice and recovered both times by reasoning from dates and titles.

Both finished both tasks. The difference between 97.5 and 27.5 is entirely the cost of finding the way, which is precisely what §9's findings address.

8. Probe question responses

Asked after the SUS, same wording every session (Vietnamese wording in

`docs/usability_testing/design/SUS_Instrument_VI_EN.md` §4). Summarise across the five participants and quote verbatim where a quote carries the point; attribute each quote to P1–P5 and to a timestamp in that session's notes.

Probe	What the five said	Quotes (attributed, verbatim)	Feeds finding
Clarity — unsure what the platform wanted, or what would happen next	Splits cleanly in two. P3 understood the form immediately. The other four each name a specific moment of doubt, and they are not the same moment: where to start (P2, P4), which request type to pick (P1, P5), and which row in the list is mine (P4).	P4: "Tôi không chắc nên tìm ở Hồ sơ, Sự kiện hay Hỗ trợ, và ở danh sách tôi cũng không biết dòng nào là yêu cầu vừa gửi." · P2: "Ban đầu tôi tìm trong trang Sự kiện vì nghĩ vấn đề đăng ký sẽ nằm ở đó; tôi không đoán ngay phải vào hỗ trợ." · P1: "Ở phần chọn loại yêu cầu tôi dừng lại một chút vì chưa chắc trường hợp mất đăng ký thuộc mục nào." · P5: "tôi phân vân giữa hai loại yêu cầu và không chắc đổi lựa chọn có làm mất nội dung đã nhập không." · P3: "Tôi hiểu ngay biểu mẫu cần gì."	D-024, D-025, D-026
Error recovery — noticing a mistake and getting back on track	Recovery worked, and worked <i>because of validation</i> , not because of navigation. Both people who erred on T1 were caught by the form refusing to submit and fixed it unaided. Recovery on T2 was worse: it meant re-reading list rows.	P2: "Khi bấm gửi mà ảnh chưa có, phần báo lỗi làm tôi nhận ra và quay lại gắn ảnh." · P4: "Tôi nhận ra mô tả chưa đủ khi biểu mẫu không cho gửi; về sau tôi mở nhầm yêu cầu cũ rồi dựa vào ngày và tiêu đề để quay lại." · P1: "khi quay lại danh sách, dòng mới nhất và trạng thái giúp tôi biết mình đang đúng chỗ."	D-025 (T2 half); the T1 half is a strength, see §10
Speed — anything slower or more effortful than it should be	Nobody said the <i>system</i> was slow. Every complaint is about effort spent searching, not waiting.	P4: "Tốn công nhất là dò nhiều menu và đọc từng yêu cầu để tìm đúng cái mới." · P2: "Tìm chỗ báo vấn đề và tìm lại yêu cầu hơi lâu vì tên menu chưa giống cách tôi nghĩ." · P5: "tôi phải kiểm tra các trường và làm mới trang để chắc phản hồi đã cập nhật." · P3: "việc đính ảnh và gửi diễn ra nhanh."	D-024, D-025
Trust — confidence that the report went through and someone would see it	All five ended up confident, but only two were confident <i>at the moment of submitting</i> . The other three needed to go and look at the list or the detail page first — trust came from persistent state, not from the confirmation.	P4: "Tôi chưa thật sự tin ngay sau khi gửi vì thông báo biến mất nhanh; chỉ khi mở được chi tiết và thấy trạng thái hoặc phản hồi tôi mới yên tâm." · P2: "Tôi tin ở mức vừa; có thông báo đã gửi nhưng chỉ yên tâm hơn khi thấy yêu cầu nằm trong danh sách." · P5: "Tôi khá tin vì yêu cầu có mã và trạng thái trong danh sách." · P3: "Tôi rất tin vì có xác nhận gửi, mục yêu cầu lưu lại nội dung, và thông báo dẫn tới phản hồi chính thức."	D-027

9. Findings, ranked by severity

Severity scale (Nielsen 0–4, as required by §6 Task 2 Phase 3):

0	Not a usability problem
1	Cosmetic — fix if time permits
2	Minor
3	Major — high priority
4	Catastrophe — fix before release

Rules applied when turning observations into the rows below, per

`docs/usability_testing/00_Run_Plan.md` §4.3:

- Group by **cause, not symptom** — three participants stumbling in three places for one underlying reason is **one** finding.
- Systemic** (a finding): hit in **≥ 2 of 5** sessions, *or* once with a structural cause that can be named.
Isolated slip (not a finding): one participant, once, no structural explanation — those go in §10, not

here.

- Every row carries a **screenshot**; the brief asks for one per ranked finding by name. Files live in `reports/evidence_task2/`; an inline image embed from this file is written `![alt]` (`../reports/evidence_task2/<file>`).
- Distinguish a **product defect** (goes to the findings log as a Bug) from a **design problem** (goes as Usability) in the Type column.

#	Severity	Type	Finding	Affected (of 5)	Evidence	Root cause	Recommendation	Log ID
F1	3 — Major	Usability	Nothing on the platform tells a user that "I have a problem with my registration" is handled under Support. Two of five began the task somewhere else entirely, and neither reached the support form by recognition — they reached it by elimination after the pages they expected turned out not to offer it.	2/5 (P2, P4)	P2: "Ban đầu tôi tìm trong trang Sự kiện vì nghĩ vấn đề đăng ký sẽ nằm ở đó; tôi không đoán ngay phải vào hỗ trợ." · P4: "Tôi không chắc nên tìm ở Hồ sơ, Sự kiện hay Hỗ trợ", "Tốn công nhất là dò nhiều menu" · Recordings Session_P2.mp4 , Session_P4.mp4 . No single-frame still: this finding is about movement across pages, which one frame cannot show.	The entry point is named after the organisation's internal category ("Hỗ trợ") rather than after the user's situation. A registration problem is mentally filed under the event, and the event page offers no route onward.	On the event / registration page, add a visible "Bảo vấn đề về đăng ký này" action that deep-links into /complaints/new with the event pre-filled. Failing that, at minimum surface Support from the event context.	D-024
F2	3 — Major	Usability	A user cannot tell which row in their request list is the one they just filed. The list gives no "just submitted" cue, so identifying your own newest request means reading rows and reasoning from dates and titles. One participant opened the wrong request; a second described the same search as slow.	2/5 (P4 hit it, P2 describes it)	P4: "ở danh sách tôi cũng không biết dòng nào là yêu cầu vừa gửi", "tôi mở nhầm yêu cầu cũ rồi dựa vào ngày và tiêu đề để quay lại" · P2: "tìm lại yêu cầu hơi lâu vì tên menu chưa giống cách tôi nghĩ" · Screenshot D-025_P4_two_pending_rows_same_timestamp.png — P4's own list showing two Pending requests both stamped Aug 3, 2026, 1:40 PM, with nothing marking which one he had just filed; recording Session_P4.mp4 .	After submitting, the user is returned to a list with no continuity from the action they just performed — no highlight, no anchor, no "mới nhất" marker. Default ordering alone does not communicate recency to someone who does not know the ordering.	Return the user to the detail page of the request just created , not to the list. If the list must be the landing page, highlight the new row and anchor to it.	D-025
F3	2 — Minor	Usability	The request-type choice is guesswork, and users fear that changing it will discard what they have typed. Two participants stalled at the same control, and one of them was reasoning about data loss rather than about categories.	2/5 (P1, P5)	P1: "Ở phần chọn loại yêu cầu tôi dừng lại một chút vì chưa chắc trường hợp mất đăng ký thuộc mục nào" · P5: "tôi phân vân giữa hai loại yêu cầu và không chắc đối lựa chọn có làm mất nội dung đã nhập không" · Screenshot D-026_P2_request_type_options_unexplained.png — the open dropdown, all four options as bare labels (Support · Complaint · Contact · Other), no example or description on any of them	Type options are unexplained labels with no examples, so the user cannot map their situation to one. The data-loss fear is a separate, second-order cost of that same uncertainty.	Add one line of example text per request type ("ví dụ: đã đăng ký nhưng không check-in được"). Guarantee — visibly — that switching type preserves entered content.	D-026
F4	2 — Minor	Usability	The submission confirmation is too transient to establish trust. Three of five did not believe the request had been filed until they navigated somewhere else and saw it persisted. The toast disappears before it has done its job.	3/5 (P2, P4, P5)	P4: "Tôi chưa thật sự tin ngay sau khi gửi vì thông báo biến mất nhanh; chỉ khi mở được chi tiết và thấy trạng thái hoặc phản hồi tôi mới yên tâm" · P2: "có thông báo đã gửi nhưng chỉ yên tâm hơn khi thấy yêu cầu nằm trong danh sách" · P5: "yêu cầu có mã và trạng thái trong danh sách" · Recordings Session_P4.mp4 , Session_P5.mp4 . No still: every post-submit frame in P1, P4 and P5 was searched and none contains the confirmation toast — consistent with the finding, but absence of a frame is not offered as proof of it.	A time-limited toast is the only acknowledgement of a state-changing action the user cares about. Trust is being carried by <i>persistent</i> artefacts (the ID, the status pill) that the confirmation itself never shows.	Replace the transient toast with a persistent confirmation that names the request ID and its status, on the page the user lands on after submitting. Pairs naturally with F2's fix.	D-027

Corroboration of an existing Task 1B finding. P5 tried the notification bell first and abandoned it: "*lúc xem phản hồi tôi thử chuông rồi quay lại vì không thấy đúng thông tin, sau đó tìm trong danh sách yêu cầu.*"

That is the user-side consequence of **D-015** (notification summaries carry a permanently empty complaint title), already logged from Task 1B checklist execution. It is recorded here as independent confirmation from a real user rather than as a new finding — the root cause is already logged, and merging by cause is the rule this project applies everywhere else.

10. Observations that are not findings

Single-participant slips with no structural cause, and anything interesting that does not meet the bar in §9. Recorded so the evidence is complete, explicitly **not** counted as findings and not ranked — padding the finding count with these is how a small study starts overclaiming.

Observation	Participant	Why it is not promoted to a finding
Refreshed the page to be sure the response had updated: " <i>tôi phải kiểm tra các trường và làm mới trang để chắc phản hồi đã cập nhật</i> "	P5	One participant, once. Plausibly the same root cause as F4 (trust carried by persistent state, not by feedback), but a single instance with a second possible explanation — ordinary caution — does not meet the ≥ 2 bar, and folding it into F4 would overstate F4's reach.
Reached the response via the notification bell without difficulty	P3	Not a problem at all. Recorded because it is the <i>only</i> successful bell route in five sessions and is the counterexample to P5's failed one — worth keeping when D-015 is fixed and re-tested.

A strength worth stating, since it is a real result and not padding. Form validation on D1 did its job in both cases where it mattered: P2 submitted without an image and P4 with too thin a description, and **both noticed the problem and fixed it unaided, from the form's own error message** — "*phần báo lỗi làm tôi nhận ra và quay lại gắn ảnh*" (P2), "*Tôi nhận ra mô tả chưa đủ khi biểu mẫu không cho gửi*" (P4). Error prevention and recovery on the form itself is the part of this flow that works, and no recommendation below asks for it to be changed.

11. Prioritised recommendations

Concrete and actionable — name the screen, the control, and the change. Ordered by severity first, then by cost to fix, so the top of the list is what to do on Monday morning. Each row must trace back to a finding number in §9; a recommendation with no finding behind it is an opinion.

Priority	From	Screen / control	Concrete change	Severity	Effort
1	F2 (D-025)	D1 → post-submit redirect	After a successful submit, land the user on <code>/complaints/{id}</code> — the request they just created — instead of the list. Removes the "which row is mine" problem at source rather than decorating around it.	3	S
2	F1 (D-024)	Event / registration page → D1	Add a "Báo vấn đề về đăng ký này" action on the event context, deep-linking to <code>/complaints/new</code> with the event pre-filled. Puts the entry point where users already look.	3	M
3	F4 (D-027)	D1 confirmation	Replace the transient toast with a persistent confirmation block showing the request ID and status. If priority 1 ships, this is where it lands anyway — the two fixes combine.	2	S
4	F3 (D-026)	D1 request-type select	One line of example text under each type option; guarantee and state that switching type preserves entered content.	2	S
5	D-015 (Task 1B, confirmed by P5)	D5 notification summary	Populate the complaint title in notification summaries. Currently the bell is a dead end for anyone who tries it — P3 succeeded through it only because the surrounding cues were enough.	—	S

Priorities 1 and 3 are the same change. Landing the user on the detail page of the request they just created both identifies the request (F2) and provides a persistent confirmation (F4). One small fix addresses a Major and a Minor finding, which is why it is first despite F1 also being Major.

11b. Conformance against §6 Task 2, item by item

Stated as a table so nothing has to be inferred. Every row names where the requirement is met, or says plainly that it is not and why it cannot be recreated after the fact.

§6 requirement	State	Where / why
Phase 1 — goal-oriented task scenario, no click path	Met	§3; wording avoids EMS's own vocabulary so the task tests findability
Measure: task success	Met	§6, per participant in <code>results/Metrics_Table.md</code>
Measure: time on task	Not met	No clock was run during the sessions. The recordings are short excerpts (18.8–26.0 s of captured footage), not full sessions, so a duration read off them would not be time on task. Cannot be recreated without re-running the study.
Measure: error / hesitation count	Not met	No observation log was kept. Self-caught errors are recorded where a participant described one (2 on T1), but that is their account, not a count. Same reason as above.
Measure: post-task SUS or UEQ-S	Met	SUS, 10 items, §7; scored by <code>score_sus.py</code>
Open probes covering clarity · error recovery · speed · trust	Met	§8, all four asked of all five, same wording
5 real participants , target profile, verifiable contacts, middle four digits masked, outside this class	Met	§5; four universities, none in this course, none in software
Pilot with one extra person	Not met	No pilot ran. Instruments were desk-checked instead. Cannot be added after the counted sessions without inventing a participant.
Phase 2 — "testing the product, not you" + think-aloud framing	Met	<code>design/Moderator_Runsheet.md</code> §5, spoken verbatim each session
Observe neutrally, no leading hints	Met	Do-not-say list, same file
Record the screen (and audio, with consent)	Met	Five distinct recordings, <code>reports/evidence_task2/Session_P1..P5.mp4</code> ; verified 0 % frame overlap between every pair
Structured notes on friction, errors, hesitations, verbalised frustration	Partly met	What exists is each participant's own written account of those things (<code>results/session_notes/</code>), not a moderator's structured log taken during the session. Honest label: participant-reported, not observer-recorded.
Close each session with the scale, then the probes	Met	§1; SUS administered before the probes, deliberately
Phase 3 — score the scale across the five	Met	§7, all five individual scores plus mean, SD and range
Tabulate task metrics (success rate, mean time, errors)	Partly met	Success rate yes; mean time and errors are the two measures above that were never taken
Group similar pain points; separate isolated bugs from systemic design issues	Met	§9 groups by cause; §10 holds what did not meet the systemic bar and says why

§6 requirement	State	Where / why
Rank findings by severity 0–4	Met	§9, Nielsen scale, two at 3 and two at 2
Report: scenario · participant table · metrics table · ranked findings · recommendations	Met	§3 · §5 · §6–§7 · §9 · §11
A screenshot per ranked finding	Partly met	D-025 and D-026 carry a still cut from the recording of the participant who hit them. D-024 is about movement across pages and D-027 about a toast that has already vanished — neither is a single frame; both cite the recordings instead, and each says so in its own row rather than pointing at an unrelated image.
Log genuine findings through the §7 channel	Outstanding	D-024...D-027 are in <code>docs/05_Bug_Usability_Findings_Log.md</code> and not yet submitted to the Google Form — see §13

Three requirements cannot be closed retroactively: time on task, error/hesitation counts, and the pilot. Each needed something to happen *during* the sessions, and none can be supplied afterwards without inventing data — which §12 makes grounds for voiding the whole task. They are reported as unmet rather than filled in.

12. Limitations

- **n = 5** supports discovery, not measurement. Five was set by the brief, not chosen from a cost–benefit curve; no percentage, confidence interval or significance claim is made anywhere in this report.
- **Scope is narrower than Task 1B.** Participants exercise D1 and D2 (and D5 only if they went there themselves). **D3/D4 are driven by the moderator**, because a usability participant cannot be handed an admin account — so the admin side has checklist coverage (Task 1B) but no user-testing coverage.
- **Single environment.** One browser, one device class, one network, per §1 — cross-platform behaviour is Task 3's evidence (`docs/04_Task3_Cross_Platform_Matrix.md`), not this report's.
- **Recruiting bias.** Participants came from the moderator's own network (see `docs/usability_testing/design/Recruiting_Kit.md` §2), which skews younger and more digitally literate than the full EMS population; findability results are therefore, if anything, optimistic.
- **The moderator is also the report's author**, so the observation and the analysis are not independent. Mitigated by fixing the success criteria and the probe wording before the first session, not after seeing the data.
- **Known defects were present during the sessions** (`docs/usability_testing/design/Moderator_Runsheet.md` §4). Where one caused a task failure, the failure is the product's and is recorded as such — but it also means task times are not a clean measure of the intended design.
- **No pilot session was run.** §6 Task 2 Phase 1 asks for one before the counted sessions. The instruments were desk-checked against the brief instead, with task wording, success criteria and probe wording all fixed in advance — but nothing was rehearsed against a real person, so any wording problem was paid for by the counted five rather than by a throwaway session.
- **Time on task, error counts and hesitation counts were not measured.** No clock was run and no observation log was kept. These are named in §6's minimum measure set and this study does not carry them; the remaining measures (task success, SUS, probes) are complete for all five.
- **Task outcomes are self-reported, not moderator-observed.** They are reconstructed from what each participant wrote, sentence by sentence, in `docs/usability_testing/results/Metrics_Table.md`. The

five recordings in `reports/evidence_task2/` can confirm them and have not yet been reviewed for that purpose.

- **All five sessions ran on one day (2026-08-03)**, so no fix could be trialled between sessions and every participant met the same build.

13. Handoff — findings channel (§7) and evidence

Every genuine defect and every usability improvement from these sessions is logged in `docs/05_Bug_Usability_Findings_Log.md` from **D-024 onwards** (D-001...D-023 are taken — D-020...D-022 came from Task 3, and **D-023 was allocated by Task 1B on 2026-08-02**, after this section was first written; D-013, D-014 and D-018 are retired and must not be reused), typed `Bug` or `Usability` with the 0–4 severity, and **each one also submitted to the Google Form** named in §7 of the brief. The log and the form must agree — the TA may cross-check the counts.

Finding # (§9)	Findings-log ID	Severity	Submitted to the form
F1	D-024	Usability 3	Not yet submitted
F2	D-025	Usability 3	Not yet submitted
F3	D-026	Usability 2	Not yet submitted
F4	D-027	Usability 2	Not yet submitted

Together with **D-023** (raised by Task 1B on 2026-08-02, also unsubmitted), that is **five form entries outstanding**; the log will then hold 24 findings and the form 24. §7 asks the two to agree and the TA may cross-check the counts.

Evidence index. Session recordings and screenshots are in `reports/evidence_task2/`, named `P<n>_<task>_<what-it-shows>.<ext>` per `docs/usability_testing/design/Moderator_Runsheet.md` §7. Raw session notes are in `docs/usability_testing/results/session_notes/`; raw SUS answers in `docs/usability_testing/results/SUS_Responses.csv`.

AI assistance on this report (study design, scoring, clustering, drafting) is declared in `docs/06_AI_Audit_Report.md`, which states explicitly that **the sessions, the participants and the session data were not AI-produced**.