

AI Critique — HW04 (Mandatory, \$10)

Student: Lê Phạm Kiều Duyên · **Student ID:** 23127184

Word count: 263 words (essay body)

The AI accelerated the initial Playwright structure, but several confident outputs were wrong or incomplete. Its first configuration pointed `baseURL` to port 3000, which is the EShop API rather than the customer web application. It also generated `getByLabel` locators without checking whether the visible labels were programmatically associated with their inputs, and guessed button text that did not exist. More seriously, in FR-11 it derived some expected values from the current implementation instead of the requirement. This produced contradictory cancellation oracles: one row expected a shipping order to expose a cancel action while other rows correctly required shipping cancellation to be refused. A defensive optional-field guard then allowed one control case to skip its data-integrity assertion entirely.

These failures occurred because the first generation pass treated documentation and common UI conventions as sufficient evidence. Cases were considered independently, so the model did not compare expectations across the whole dataset. Runtime probing helped verify selectors, but it became dangerous when observed product behaviour was reused as the oracle. The AI also preferred fail-soft code that avoided exceptions, even when a missing field should have failed fast as invalid test data.

I learned that effective AI collaboration requires separate evidence gates. Requirements define expected behaviour; source and DOM inspection verify how to reach it; execution reveals whether the product conforms. These sources must not be substituted for one another. I now require selector-count probes, cross-row consistency checks, fail-fast data validation, and a human review stop before execution. A red assertion is preserved until it is classified, rather than weakened to make the suite appear successful.