

HW05 - AI-assisted Performance Testing

Student ID: 23127184 System under test: EShop REST backend Date of official execution: 2026-08-17 (Asia/Saigon)
Workflow: Login → Search product → Product detail → Add to cart → Checkout

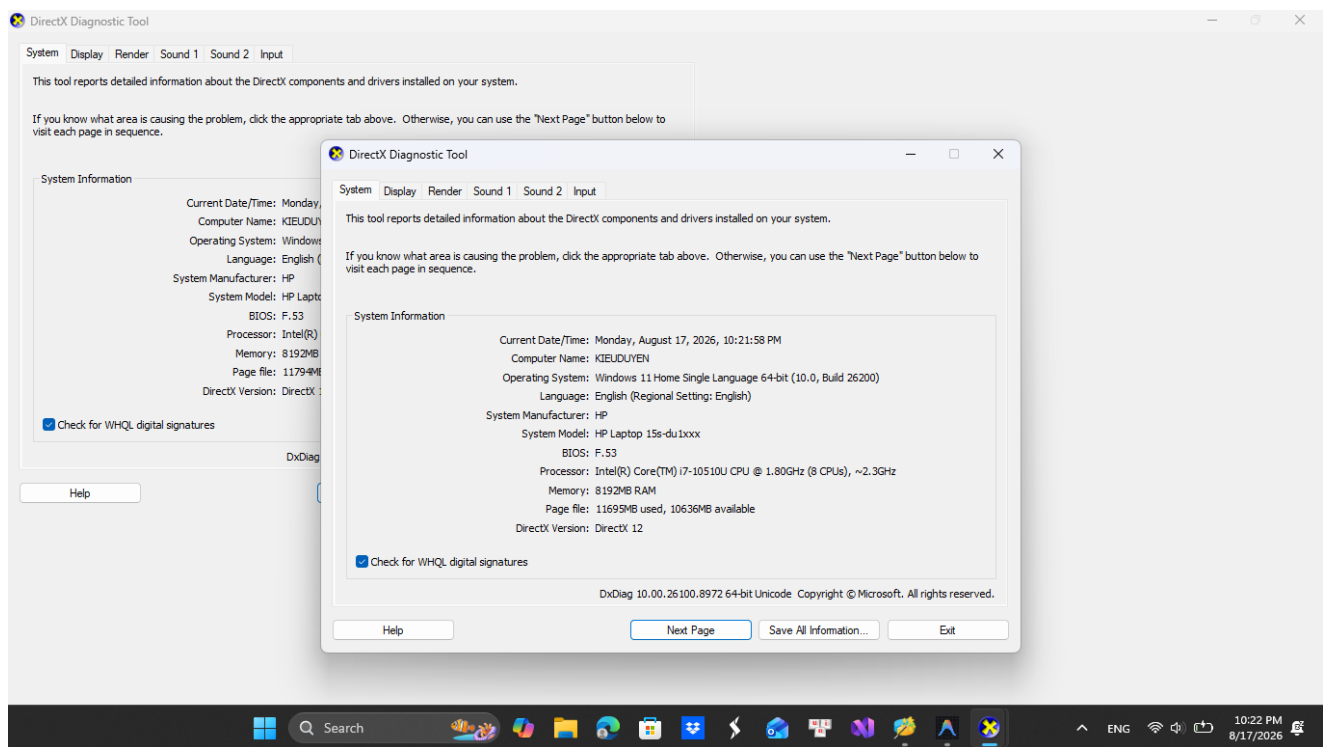
1. Executive summary

This assignment applies the seven-step performance-testing method to one data-driven end-to-end purchase journey covering auth-heavy, read-heavy and transactional endpoint groups. The student confirmed that the workflow is not duplicated by another group member. OpenAI Codex assisted with requirements review, plan generation, execution automation and initial analysis; all reported measurements come from genuine local JMeter runs and were reviewed against the raw JTL files.

Load, Stress, Spike and a 15-minute Soak completed without request or assertion failures. Load sustained 64.16 req/s overall at request p95 10 ms. Stress scaled to about 260.6 req/s at 132 VU with p95 13 ms, so the tested range did not reach the knee. Spike reached 361.6 req/s at p95 26 ms and recovered to baseline within 30 seconds. Soak sustained 52.8-54.5 req/s at p95 no higher than 12 ms in full stable buckets. The demonstrated endurance threshold is at least 54.5 req/s (about 10.9 complete journeys/s) with a 172.0 MB backend working-set ceiling on this laptop.

2. Environment and evidence identity

Item	Value
Hostname	KIEUDUYEN
CPU	Intel Core i7-10510U, 4 cores / 8 logical processors
RAM	8,192 MB physical
Operating system	Windows 11 Home Single Language, build 26200
Java / JMeter	Java 21.0.12 / Apache JMeter 5.6.3
Node.js	24.4.1
Backend	http://localhost:3000, rate limiter bypassed with LOADTEST=1
Seeded state per run	242 users total, 5 products, 0 orders before traffic



Hardware identity and specification

JMeter and the SUT shared the same 8 GB host. Before every official scenario, the backend was stopped, restarted to recreate SQLite and clear in-memory carts, 240 accounts were registered through the API, all CSVs were validated and the new backend PID was passed explicitly to the resource sampler.

3. Criteria fixed before execution

Scope	p95 limit	Error/pass rule
Login	250 ms	Error <1%; token present
Search and detail	200 ms	Error <1%; content assertion passes
Cart and checkout	300 ms	Error <1%; order ID returned
Complete journey	1,000 ms	At least 99% endpoint-complete journeys pass

Spike also had to return to within 20% of pre-burst p95 within 120 seconds. Soak had to hold 10-15 minutes, retain the latency/error criteria and report a measured RPS and memory ceiling. Criteria were not relaxed after results were known.

4. AI-assisted plan and human corrections

All official plans share the same correlated workflow and three CSV pools: credentials, live product-derived search/ID values and checkout payloads. Login extracts a JWT; downstream requests use it. Assertions check content as well as HTTP status. Load uses Summary Report, Stress uses Aggregate Report and Spike uses error-only View Results Tree, satisfying the distinct-view requirement.

Human review corrected the AI-generated approach in several material ways:

- inherited round thread counts were replaced with values derived from the measured 30.4 ms service time plus disclosed compressed think time;
- Stress was changed from one linear ramp into four measurable accumulating stages;
- absolute CSV paths were replaced by portable `data.dir` properties;

- a failing-login account was excluded from the successful journey pool to prevent lockout contamination;
- weak status-only checks were replaced by token, product-ID and order-ID assertions;
- the runner was changed to require an explicit backend PID; and
- controller rows were no longer trusted when the scheduler cut off an iteration before all five endpoints completed.

5. Scenario design

Scenario	Executable profile	Question
Load	34 VU, 68 s ramp, 668 s total	Can anticipated peak load remain healthy?
Stress	33 → 66 → 99 → 132 VU, 600 s	Where does scaling stop or an SLO fail?
Spike	17 baseline + 168 burst, 270 s	Can a sudden 185-VU peak recover?
Soak	27 VU, 54 s ramp, 954 s total	Is 15-minute sustained behavior stable?


Each HTTP sampler uses a Uniform Random Timer of 300-700 ms. This is compressed relative to human browsing because hundreds more co-located threads would make the laptop/JMeter the test subject. The report therefore describes an arrival rate model, not literal shopper reading time.

6. Results

Metric	Load	Stress	Spike	Soak
Request samples	42,810	94,104	28,325	49,727
Complete journeys	8,547	18,775	5,591	9,934
Requests/s overall	64.16	157.00	105.26	52.17
Complete journeys/s	12.83	31.39	20.91	10.44
Mean response time	4.0 ms	4.3 ms	6.3 ms	4.2 ms
Average latency	4.02 ms	4.30 ms	6.27 ms	4.17 ms
Request p95 / p99 / max	10 / 20 / 52 ms	11 / 19 / 74 ms	19 / 37 / 111 ms	11 / 16 / 40 ms
E2E p95 / max	34 / 77 ms	36 / 88 ms	64 / 192 ms	31 / 73 ms
Error rate	0.00%	0.00%	0.00%	0.00%
Bandwidth	27.4 KB/s	67.2 KB/s	45.1 KB/s	22.3 KB/s
Backend CPU mean / peak	2.2 / 4.4%	4.4 / 8.2%	2.9 / 11.1%	2.0 / 3.7%
Memory start / peak / end	52.2 / 172.2 / 73.7 MB	52.8 / 172.8 / 122.2 MB	53.7 / 178.4 / 64.3 MB	52.8 / 172.0 / 89.6 MB

6.1 Load

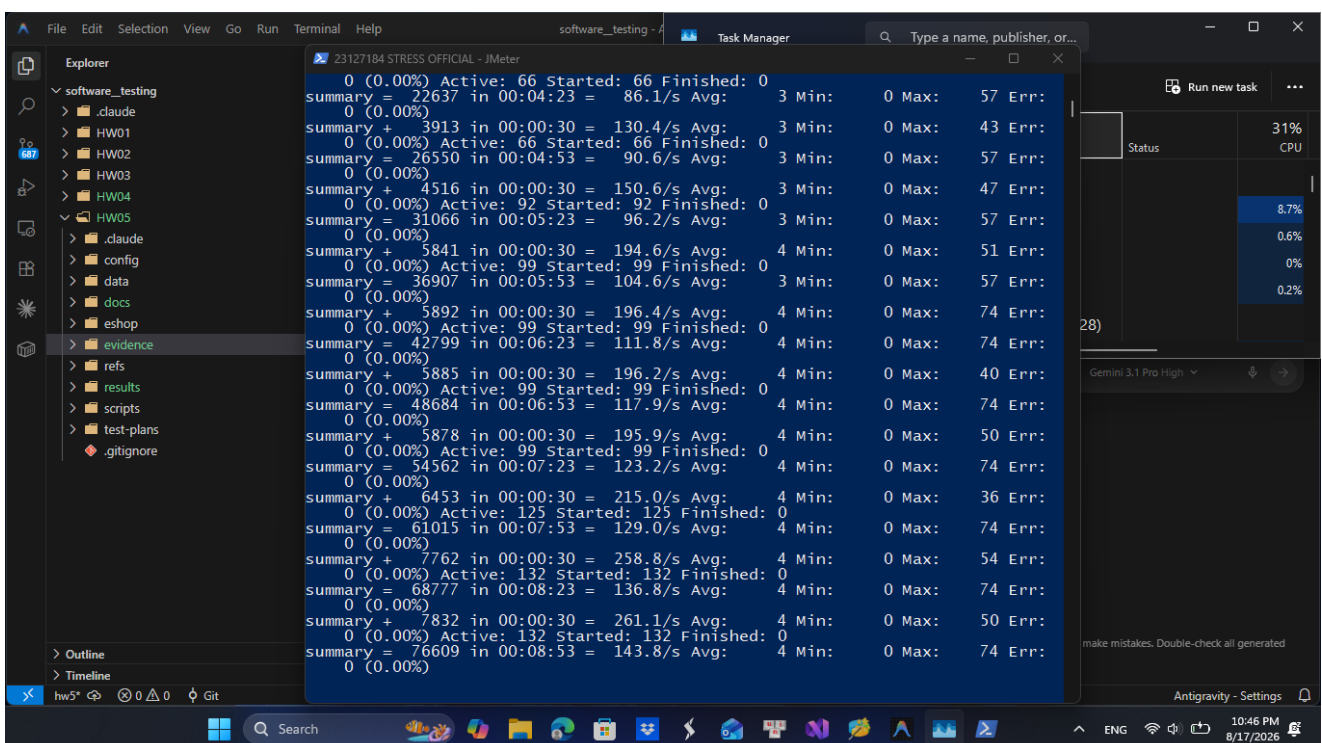
Load delivered 12.83 complete journeys/s including ramp, 97.5% of the planned 13.16 journeys/s; stable request buckets were about 67.3 req/s. All endpoint and E2E thresholds passed.



Load stable evidence

6.2 Stress

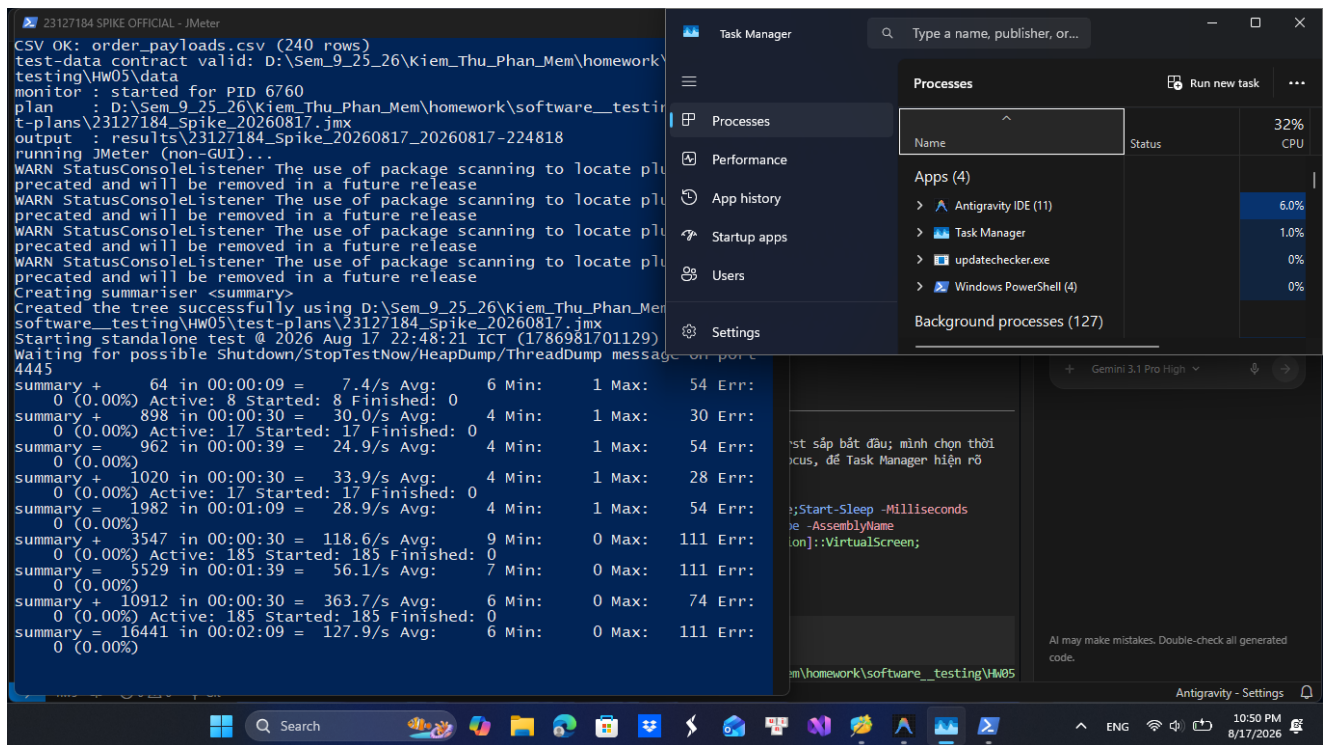
Stable stages produced about 65.9, 130.7, 196.1 and 260.6 req/s. Their request p95 stayed between 10 and 13 ms with zero errors. Throughput continued to scale, so 132 VU is not called the capacity. The defensible finding is that the knee is above the tested range under this configuration.



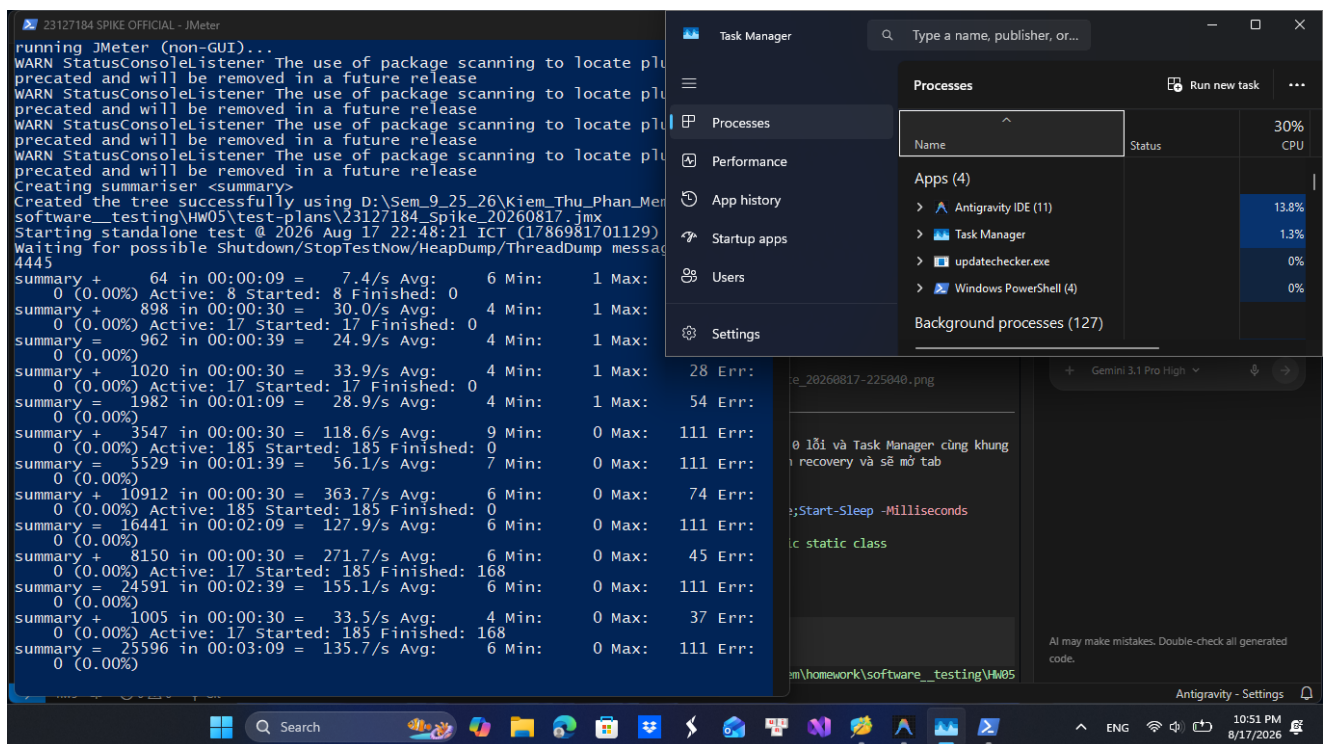
Stress peak evidence

6.3 Spike

The pre-burst phase was about 34.1 req/s at p95 11 ms. The burst reached 361.6 req/s at p95 26 ms. In the first full recovery bucket, throughput returned to 33.8 req/s and p95 12 ms, only 9.1% above baseline and inside the 20% rule.



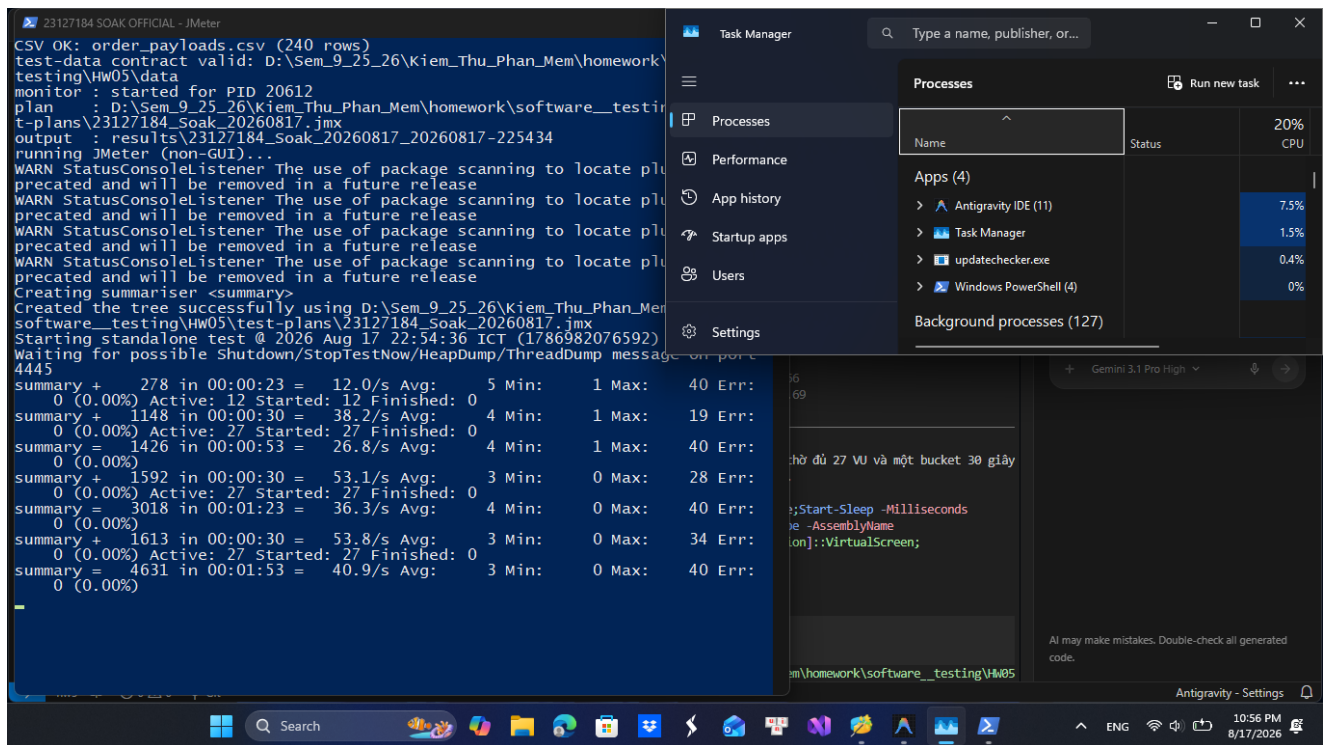
Spike burst evidence




Spike recovery evidence

6.4 Soak and endurance threshold

Twenty-seven VU held for 900 seconds after ramp. Full stable buckets stayed between 52.8 and 54.5 req/s, p95 stayed at or below 12 ms and errors stayed at zero. The backend working set started at 52.8 MB, peaked at 172.0 MB and ended at 89.6 MB. A mid-run fall from about 170 MB to 68 MB shows garbage collection, so this is not proof of a monotonic leak. End memory remains elevated and the source shows carts are never cleared after checkout; a longer recovery test is warranted.



Soak stable start



Soak late stable

7. AI analysis misinterpretation hunt

The unchanged generated analyses are stored in docs/analysis/; the reviewed corrections are in docs/reviews/03_ai_analysis_review.md.

1. One 300 ms analyzer threshold could not represent the four pre-registered label/E2E budgets. Human review applied each threshold separately.
2. Successful controller rows overcounted complete journeys. For Stress, 18,797 controller rows were marked successful but only 18,775 journeys contained all five endpoint labels.

3. The highest Stress stage was not a capacity measurement. At 132 VU, throughput still scaled and p95 remained 13 ms.
4. Overall Spike p95 19 ms hid phase behavior; the relevant values were 11 ms baseline, 26 ms burst and 12 ms recovery.
5. Soak end memory above start was only an accumulation signal. A large mid-run drop contradicted an automatic leak conclusion.


8. Optimization judgement

Clearing `userCarts[userId]` after successful checkout is supported by both code and the workload model. Indexing `users.email` and enabling SQLite WAL are plausible but unverified because no login/write bottleneck was measured. A generic network connection pool is unfounded because this SUT uses one shared `sqlite3.Database` handle. A conventional B-tree index for `LIKE '%keyword%'` is also unfounded because the leading wildcard prevents a prefix lookup.

9. Continuous performance testing

The proposed pipeline runs a short Load gate only for performance-sensitive commits, confirms regressions before blocking, schedules Stress/Spike and runs Soak for release candidates. It compares each label and the complete journey with the median of ten same-runner baselines. It blocks on an absolute SLO or a reproduced p95 increase greater than both 20% and 10 ms.

Commit-aware performance regression gate



Scheduled lane: nightly Stress + Spike; release-candidate Soak

The trade-off is cost versus confidence: per-commit Soak would be slow and expensive, while one noisy Load result would create false alarms. Path filters, a fixed runner, an absolute 10 ms noise floor and confirmation runs reduce cost and false positives; scheduled high-load tests reduce false negatives.

10. AI Critique (272 words)

The AI was useful for generating a repeatable JMeter workflow, but several of its early conclusions were too confident. First, it treated the single-user baseline capacity of 32.9 journeys per second as if it predicted multi-user capacity. The Stress run later reached about 52 complete journeys per second at 132 VU without a knee, showing that the baseline value was only a sizing heuristic. Second, the initial analyzer trusted successful Transaction Controller rows. Manual comparison of the five endpoint-label counts showed that scheduler shutdown produced partial journeys: for example, Stress had 18,797 successful controller rows but only 18,775 endpoint-complete journeys. The analyzer was corrected to report the conservative lower bound. Third, an automatic warning described end memory above start as a possible accumulation signal. That warning was reasonable, but calling it a leak would have been wrong: the Soak trace fell from about 170 MB to 68 MB during garbage collection and ended at 89.6 MB. Source inspection instead found a specific risk: checkout does not clear the in-memory cart. Finally, generic recommendations such as a connection pool or a B-tree index for `LIKE '%keyword%'` sounded plausible but did not match this implementation. It uses one shared SQLite handle, and a leading wildcard prevents a normal prefix index from serving the query. These misses occurred because the model initially reasoned from common performance patterns before it had raw label counts, phase-specific timelines, and source code. I learned to treat AI output as a testable hypothesis: pre-register criteria, preserve raw evidence, recompute every important number, inspect the actual code path, and make capacity or leak claims only within the measured scope.

11. Conclusion and limitations

All four executed scenarios met their pre-registered latency, assertion and error criteria. No genuine error response, crash or performance regression was observed, so no GitHub performance issue was fabricated. The results apply only to a five-product local database, compressed think time and a co-located load generator with the rate limiter bypassed. They do not establish production/WAN capacity. A 198-VU Stress retest is prepared but should run only after the host again has at least 2 GB free RAM.

Demo video and Agent Skill demonstration: VIDEO_URL_PENDING. The remaining anti-cheat submission action is for the student to record/upload an unlisted video of at least six minutes with their own Vietnamese narration, then replace this marker and the two markers in README.md before packaging.