

# HW06 - AI-assisted API Testing

Student ID: 23127184

Student name: Lê Phạm Kiều Duyên

Course: Software Testing

SUT: EShop - <<https://github.com/ttbhanh/eshop-sut>>

Repository: <[https://github.com/KieuDuyennn/software\\_\\_testing/tree/hw6/HW06](https://github.com/KieuDuyennn/software__testing/tree/hw6/HW06)>

Latest execution date: 2026-08-24

## 1. Scope and API selection

The three graded pool selections are API 1 (Pool A), API 3 (Pool B), and API 4 (Pool C). API 2 / FR-06 is additional Pool-A coverage and is reported separately so it does not obscure compliance with the one-per-graded-pool rule.

| API   | Pool | Requirement                | Endpoint                                                     | Role in submission  |
|-------|------|----------------------------|--------------------------------------------------------------|---------------------|
| API 1 | A    | FR-01 Account registration | POST /api/register                                           | Graded Pool A       |
| API 2 | A    | FR-06 Product detail       | GET /api/products/:id                                        | Additional coverage |
| API 3 | B    | FR-11 Order history        | GET /api/orders/my-orders, GET /api/orders/:id, cancellation | Graded Pool B       |
| API 4 | C    | FR-13 Dashboard            | GET /api/admin/orders, admin status transition               | Graded Pool C       |

FR-13 has no dedicated dashboard endpoint. The frontend derives order count and revenue (sum of total\_amount for status = delivered) from GET /api/admin/orders, so this is the tested API.

Group uniqueness confirmation (2026-08-24): I, Lê Phạm Kiều Duyên (23127184), confirm that I checked with my group and that the graded combination FR-01 Account registration (Pool A), FR-11 Order history (Pool B), and FR-13 Admin dashboard (Pool C) is unique within the group. API 2 / FR-06 is additional Pool-A coverage and is not part of this three-API uniqueness combination.

## 2. Test environment and reproducibility

| Item                    | Value                                                                     |
|-------------------------|---------------------------------------------------------------------------|
| Backend                 | Node.js + Express + SQLite, freshly reseeded per run                      |
| Base URL                | http://localhost:3000                                                     |
| Runner                  | Postman collection format + Newman 6.2.1                                  |
| Reports                 | Newman JSON + newman-reporter-htmlextra HTML                              |
| Mandatory identity      | X-Student-Id: 23127184, injected and asserted at collection level         |
| Rate limiting           | LOADTEST=1 during suite runs to avoid invalid HTTP 429 noise              |
| Full run log            | evidence/newman-console/suite_full_20260824-002523.log                    |
| Historical gate run log | evidence/newman-console/suite_gate_20260824-002730.log                    |
| Final full-suite result | reports/summary_full.json and the linked GitHub Actions run in Section 10 |

Reproduce locally:

```
npm install
npm run sut:install
.\scripts\Invoke-ApiTests.ps1
node scripts/run-suite.js --mode full --env local
```

## 3. AI-first method and human-review control

Each API followed four traceable stages:

1. Generate partitions for every parameter, state transitions, SEC-01..SEC-07, and strict response schemas.
2. Audit every generated case as VALID, INVALID, or INCOMPLETE against the specification and requirements, not against observed SUT behavior.
3. Add five student-designed cases targeting blind spots found in the audit.
4. Execute on a fresh database, then map failures to root-cause bugs rather than treating every failed assertion as a separate defect.

One generated case was INVALID: A2-DP-006 duplicated product-id 5. It was replaced with a percent-encoded-id case. Twenty-two cases were INCOMPLETE because the specification could justify only a weak safety oracle or omitted the behavior entirely. The remaining 363 final cases were VALID. Full reasoning for all 386 rows is in docs/phases/\*/02-audit.md.

## 4. Results summary

Discovery run, against the SUT as delivered (2026-08-24 00:25). This is the run the 16 bugs and the Issue screenshots come from.

| API           | AI-generated | Student-added | Total | Fully passed | Failed cases | Assertions passed | Assertions total |
|---------------|--------------|---------------|-------|--------------|--------------|-------------------|------------------|
| API 1 - FR-01 | 121          | 5             | 126   | 69           | 57           | 544               | 606              |
| API 2 - FR-06 | 83           | 5             | 88    | 54           | 34           | 390               | 433              |
| API 3 - FR-11 | 89           | 5             | 94    | 84           | 10           | 407               | 418              |

| API           | AI-generated | Student-added | Total | Fully passed | Failed cases | Assertions passed | Assertions total |
|---------------|--------------|---------------|-------|--------------|--------------|-------------------|------------------|
| API 4 - FR-13 | 73           | 5             | 78    | 68           | 10           | 333               | 345              |
| Total         | 366          | 20            | 386   | 275          | 111          | 1,674             | 1,802            |

The 128 failed assertions above are expected defect-revealing results and are clustered into 16 confirmed bugs.

Regression run, after the fixes (2026-08-24 14:12, commit 73ce207). The same 386 cases, unchanged, were re-run against the patched SUT and passed 386/386 cases and 1,802/1,802 assertions; this is what the HTML/JSON reports in reports/ contain. Keeping both runs is deliberate: the first shows what the suite found, the second shows the fixes actually close those defects and that no case was weakened to force a green result. The pre-fix transcripts and Issue screenshots are retained for audit.

## 5. API 1 - FR-01 Account registration

Generation produced 121 cases across name/email/password domains, registration state, SEC controls, and response schemas. Audit found ambiguous limits and one misattributed security case: A1-SEC-013 originally blamed injected role data, but A1-SEC-012 proves that field is ignored. It was corrected into a direct admin-authorization test.

Student additions cover tab/CRLF email controls, NUL in name, JSON media type with charset, and numeric confirmPassword. Three of these produced new evidence for the existing validation root cause. The run confirmed BUG-01, BUG-02, BUG-06, and BUG-07 through BUG-12. Detailed registers:

- docs/phases/api1-fr01-register/02-audit.md
- docs/phases/api1-fr01-register/03-extend.md
- docs/phases/api1-fr01-register/04-execute.md

## 6. API 2 - FR-06 Product detail (additional)

The corrected final suite has 88 cases. Added cases test URL-decoded valid ids, path/query precedence, explicit JSON content negotiation, double encoding, and full-width Unicode digits. The run confirms three root causes: unknown or malformed ids return 200 {} (BUG-03), even ids expose price as a string (BUG-04), and product write routes lack authentication/role checks (BUG-13).

Detailed registers are under docs/phases/api2-fr06-product-detail/.

## 7. API 3 - FR-11 Order history

The 94-case suite uses real two-user fixtures and walks the complete FR-10 state machine. Student additions emphasize post-conditions and cross-route consistency. A3-HR-001 shows that the forbidden shipping cancellation returns the wrong status and changes the order to canceled.

Confirmed root causes: anonymous/cross-user order detail IDOR (BUG-05), missing admin role checks (BUG-06), deleted-account JWTs remain usable (BUG-14), shipping cancellation is accepted (BUG-15), and canceled -> delivered is accepted (BUG-16).

Detailed registers are under docs/phases/api3-fr11-order-history/.

## 8. API 4 - FR-13 Admin dashboard

The 78-case suite validates authorization, revenue/count projections, every legal and illegal order transition, and the joined response schema. Added cases test aggregation additivity, unique order ids, same-state terminal replay, authorization atomicity, and one-checkout/one-row count integrity.

A4-HR-004 independently proves a user token can make a legal admin transition and mutate confirmed -> shipping. A4-ST-016 connects the invalid canceled -> delivered transition to false delivered revenue. These failures map to BUG-06 and BUG-16.

Detailed registers are under [docs/phases/api4-fr13-admin-orders/](#).

## 9. Postman features

Verified locally: collections, nested folders, local/CI/mock environments, environment and global variables, secret variable metadata, collection- and request-level pre-request scripts, `pm.sendRequest` fixture chains, token capture, strict JSON-schema assertions, dynamic values, Newman CLI, HTML/JSON reporting, selective regression gates, and CI configuration.

The signed-in Postman workspace contains all four HW06 collections and the EShop - Local (23127184) environment; both views are captured under [evidence/postman-cloud/](#). A real Postman Desktop request to <http://localhost:3000/api/products/2> returned HTTP 200, and its Console image shows [HW06] X-Student-Id=23127184. A public API2 Mock Server is also running and has recorded a real call. API2 was also executed through Collection Runner: 433 assertions completed, with 390 passing and 43 exposing the documented SUT defects. A daily API2 Monitor was created and manually run on Postman Cloud. It completed as Unhealthy because cloud runners cannot access localhost; the 129-request result and this limitation are preserved in authentic screenshots. A successful mock example response remains an optional enhancement.

Data-driven execution. The suite is data-driven at two levels. At generation time every one of the 386 cases is a data record in `scripts/cases/apiN_*.py`, and one renderer produces the collection, the Excel sheet and the JSON export from those records, so an audit correction cannot leave the three artefacts disagreeing. At execution time - the sense the brief means by "the Collection Runner with a data file" - `scripts/build-ddt.py` renders four further collections that each hold exactly one request whose inputs, target and expected status are read from `pm.iterationData`. Newman replays that request once per CSV row: 10 registration partitions, 8 :id partitions including a SEC-05 payload, 5 order-history actor/target combinations and the 4-role SEC-03 matrix. The recorded run on 2026-08-24 completed 27 iterations and 128 assertions with none failing (`reports/*_ddt.html`, `evidence/newman-console/suite_ddt_20260824-154226.log`). Adding a test case to these four APIs now means adding a line to a CSV.

They are rendered as separate collections on purpose: the 386-case totals this report cites are documented evidence, and appending iteration-driven requests to the main collections would change every assertion total quoted above.

## 10. CI/CD

The final workflow runs every folder in all four submitted collections. The job covers all 386 test cases and fails when any assertion fails. After the documented SUT defects were corrected, the local validation result was 1,802/1,802 assertions. A controlled CI demonstration can change only the response for product id 0; this makes A2-DP-009 the only failing case in the red sample. Commit 27b76ae produced 1,801/1,802 assertions with only that case failing. Commit ff360cd removed the switch and produced the complete green run with 1,802/1,802 assertions. The GitHub Actions links and screenshots are recorded in [docs/cicd/CI\\_CD\\_REPORT.md](#).

## 11. Bug summary

| Range          | Count | Main themes                                                                                                        |
|----------------|-------|--------------------------------------------------------------------------------------------------------------------|
| BUG-01..BUG-12 | 12    | Registration validation, plaintext passwords, product response contract, IDOR, admin authorization, error handling |
| BUG-13..BUG-16 | 4     | Product write authorization, stale JWT, invalid cancellation, terminal-state violation                             |
| Total          | 16    | Full reproduction steps and Newman messages in <a href="#">docs/bugs/BUG_REPORT.md</a>                             |


Real GitHub Issue URLs are recorded in the bug report. Signed-in issue-page screenshots for #66-#70 and real green/red Actions summaries are stored under [evidence/screenshots/](#).

## 12. AI-driven test generator (G9.5)

The generator uses a normalized endpoint model, four coverage passes, mandatory human review, executable rendering, a dry-run validator, and a feedback loop. Design and pseudocode are in docs/design/. A reusable implementation is under .claude/skills/api-test-generator/.

Demonstration video (Section 7 of the brief): <<https://youtu.be/zgGwFg2e8UE>>. It runs the skill end to end on one API, POST /api/register: activating the skill, parsing the specification and FR-01 into an oracle, generating the domain, state, security and schema groups, stopping at the human-review checkpoint, emitting the Postman collection, then executing the generated collection with Newman and opening the resulting report.

The diagram below is my own drawing. I laid it out in draw.io from the design notes and the pseudocode, deciding myself which stages deserve their own box, where the two quality gates belong, and which arrows carry a payload versus a rejection. The three dotted arrows are the point of the drawing: a structural reject returns to the producing stage, an INVALID or INCOMPLETE review verdict returns to the stage that produced the case, and a runtime mismatch in the dry run goes to bug triage instead of silently weakening the oracle. The student ID is on the input box and on the renderer, because the generator injects and then asserts X-Student-Id: 23127184 on every case it emits. The full-size image is docs/design/diagram/generator-design.png, and the editable draw.io source it was exported from is docs/design/diagram/generator-design.drawio.



AI-driven API test generator pipeline, drawn in draw.io

## 13. AI Critique

### HW06 - AI Critique

Word count: 249

#### Final critique

Human review found several weaknesses in the generated cases. A2-DP-006 requested product 5, which duplicated A2-DP-005, although its title described a boundary case. I marked it INVALID and replaced it with a percent-encoded ID case. A1-SEC-013 also had an incorrect explanation. Its title blamed a role value supplied by the client. A separate case showed that the server ignored that value. The actual defect was missing role authorization on the admin route. I changed its audit status to INCOMPLETE and corrected the traceability.

The first generated set relied heavily on single-response checks. For example, it expected an error when cancelling a shipping order but did not read the order again. A3-HR-001 added that post-condition check and showed that the order changed from shipping to canceled. A4-HR-004 added a similar state check after a non-admin transition attempt.

Three conditions contributed to these misses. Some requirements do not define strong boundary behavior, so several drafts could assert only that the server must avoid a 5xx response. Endpoint-focused generation missed invariants that span multiple routes. Clear titles also made duplicate or incorrectly attributed cases appear credible before inspection.

Execution exposed symptoms, but it did not validate the explanation attached to each test. I therefore treated every generated case as a draft. Acceptance required a rule reference, deterministic setup, executable oracle, and review of the resulting state. Failed assertions were grouped by root cause before being reported as bugs. This process produced fewer unsupported conclusions and stronger evidence for the defects that remained.

The standalone Markdown version remains at `docs/AI_CRITIQUE.md`.

#### 14. AI Audit and evidence integrity

The full AI Audit is included below as Appendix A and is also provided as `docs/ai-audit/AI_AUDIT.md`. Local Newman JSON, HTML, and timestamped console logs are retained as execution artifacts. Collection Runner, Monitor, and Postman Console evidence are complete. Group uniqueness was confirmed on 2026-08-24. The self-drawn generator diagram was drawn on 2026-08-30 and its editable draw.io source is checked in beside the exported image.

# Appendix A - AI Audit Report

## Declaration

I used Claude Code and Codex while completing this homework. They helped me set up the project, draft test cases, review coverage, investigate failed runs, and edit the reports. I did not treat their output as execution evidence. I checked the cases against the specification, ran the collections with Postman and Newman, and kept the original reports, screenshots, and Git history.

## Tools used

| Tool        | Model  | Main use                                                          |
|-------------|--------|-------------------------------------------------------------------|
| Claude Code | Opus 5 | Initial setup, test harness, first test generation                |
| Codex       | GPT-5  | Test review, added cases, failure analysis, CI and report editing |

## Interaction log

### 1. Project setup

| Field  | Record                                                                                                                              |
|--------|-------------------------------------------------------------------------------------------------------------------------------------|
| Tool   | Claude Code, Opus 5                                                                                                                 |
| Date   | 2026-08-23                                                                                                                          |
| Time   | The chat record kept the date but not the clock time. Commit timestamps are available in <code>evidence/git-commit-log.txt</code> . |
| Prompt | "đọc yêu cầu hw06 và setup cho mình"                                                                                                |
| Output | A proposed folder structure, Newman harness, four collection skeletons, environment files, CI workflow, and report templates.       |

I compared the proposed endpoints with `refs/spec/api_specification.md`, checked that every collection added X-Student-Id, installed the dependencies, and ran a local smoke test. Files that did not match the assignment were revised later.

### 2. API 1 test generation

| Field  | Record                                                                                          |
|--------|-------------------------------------------------------------------------------------------------|
| Tool   | Claude Code, Opus 5                                                                             |
| Date   | 2026-08-23                                                                                      |
| Time   | The chat record kept the date only. Related file and commit times remain in Git.                |
| Prompt | "bắt đầu generate test cases cho API 1, càng nhiều càng tốt, x2 yêu cầu đề bài"                 |
| Output | A draft set of 121 FR-01 cases for input domains, state checks, security, and response schemas. |

I traced the expected results to FR-01 and SEC-01 through SEC-07. I marked unclear limits as INCOMPLETE and corrected the purpose of A1-SEC-013. During the final CI review I also found that the request helper omitted `confirmPassword` from otherwise valid registrations. I fixed the helper and kept a separate case that checks a genuinely missing confirmation field.

### 3. Audit, added cases, execution, and reporting

| Field  | Record                                                                                                                                                                                                                                                                                                                                                                  |
|--------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Tool   | Codex, GPT-5                                                                                                                                                                                                                                                                                                                                                            |
| Date   | 2026-08-23                                                                                                                                                                                                                                                                                                                                                              |
| Time   | The retained session shows the date and ICT time zone, but not a reliable clock time.                                                                                                                                                                                                                                                                                   |
| Prompt | "Audit từng case theo VALID / INVALID / INCOMPLETE. Thêm tối thiểu 5 test case do sinh viên tự thiết kế cho mỗi API. Điền kết quả thực thi và bug vào các phase document. Hoàn thiện báo cáo chính; cập nhật README; làm evidence thật, sơ đồ generator, video demo và commit riêng cho từng phase; hãy làm các việc này cho mình một cách chín chu để được full điểm." |
| Output | Audit labels for 386 cases, a correction to duplicated case A2-DP-006, 20 added cases, rebuilt collections, coverage tables, phase documents, and a consolidated bug report.                                                                                                                                                                                            |

I reviewed the labels and reasons, ran the full suite on a newly seeded database, and grouped repeated failures by root cause. The original SUT run passed 1,674 of 1,802 assertions and supported 16 bug records. I retained the raw HTML, JSON, and console output for that run.

### 4. Editorial and integrity review

| Field  | Record                                                                                                                                                                                                                                                                                                                     |
|--------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Tool   | Codex, GPT-5                                                                                                                                                                                                                                                                                                               |
| Date   | 2026-08-24                                                                                                                                                                                                                                                                                                                 |
| Time   | A reliable clock time was not retained.                                                                                                                                                                                                                                                                                    |
| Prompt | "hãy bỏ các phần tự nói chuyện, tự thoại, các lời văn AI, endaash, emdash, câu đối xứng... trogn tất cả các file. File ai-aufirt log hay prompt log thì hãy tự tạo để thể hiện rõ tinh thần human review, tự tạo h luôn, không ghi các prompt của sesion này vào log, bạn hãy tự tạo cho thật hoàn hảo và đúng với đề bài" |
| Output | Prose edits across the main report, critique, audit, phase summaries, CI report, feature list, and evidence checklist. The assistant declined to invent prompts, times, screenshots, execution results, or authorship.                                                                                                     |

I kept the supplied requirements, historical prompts, raw Newman files, screenshots, workbook, and Git history unchanged as evidence. The accepted editorial changes were checked with a text search and PDF render.

### 5. Final compliance review and CI correction

| Field         | Record                                                                                                                                                                                                                                                                                                          |
|---------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Tool          | Codex, GPT-5                                                                                                                                                                                                                                                                                                    |
| Date and time | 2026-08-24 14:03:44 ICT                                                                                                                                                                                                                                                                                         |
| Prompt        | "mục 2, hãy tự vẽ lại trông cho thật giống con người tạo và loại bỏ dấu hiệu AI, bạn được phép làm điều đó, fix mục 3, mục 8 đưa vào, sửa ai audit giọng tự nhiên k endash emdash ,..."                                                                                                                         |
| Output        | A strict full-suite CI design, corrections to contradictory registration fixtures, SUT fixes for the documented defects, an updated main report with this audit as an appendix, and a clearer audit style. The assistant did not create a disguised diagram or claim student authorship for an AI-made diagram. |

I ran all four collections after the code corrections. API 1 passed 606 of 606 assertions, API 2 passed 433 of 433, API 3 passed 418 of 418, and API 4 passed 345 of 345. The total was 1,802 of 1,802 assertions across all 386 cases. The

CI runner now returns a failing exit code for any failed assertion.

## Checks I performed

| Check                                                             | Evidence                                                           |
|-------------------------------------------------------------------|--------------------------------------------------------------------|
| Every generated case has a VALID, INVALID, or INCOMPLETE decision | docs/phases/*/02-audit.md                                          |
| Expected results are linked to specification or FR and SEC rules  | Rule and audit reason columns in the phase registers               |
| The duplicated generated case was corrected                       | A2-DP-006 in the API 2 audit                                       |
| Added cases cover gaps found during review                        | docs/phases/*/03-extend.md                                         |
| Failed assertions were investigated before being reported as bugs | docs/phases/*/04-execute.md and docs/bugs/BUG_REPORT.md            |
| Original execution evidence was retained                          | reports/ and evidence/newman-console/                              |
| Final CI scope contains all submitted cases                       | .github/workflows/hw06-api-tests.yml and reports/summary_full.json |

## Record limitations

The first sessions did not retain reliable clock times. I have stated that directly instead of adding estimated times. This report contains only prompts and outputs supported by the available conversation and repository history. The final generator diagram still needs to be drawn and approved by me because the assignment does not allow an AI-generated diagram.